

# IPAddress

---

## Contents

---

[Contents](#)

[Benefits of this Library](#)

[Code Examples](#)

[Supported IP Address Parsing Formats](#)

[Core classes](#)

[Parse String Representation of IP Address or Host Name](#)

[IP Address and Numeric Values](#)

[Networks](#)

[Prefix Length Handling](#)

[Address Sections](#)

[IP Address Ranges](#)

[Address Tries](#)

[IP Address Operations](#)

[Parse String Representations of MAC Address](#)

[MAC Address Operations](#)

[IPv6 – MAC Address Integration](#)

[Address Framework](#)

[Conversion to String Representation of Address](#)

[Searching Text of Databases for all Addresses in a Subnet](#)

[Containment and Subnet Membership](#)

[DNS Resolution and URLs](#)

[Sorting and Comparisons](#)

[Cache Classes](#)

[Make your IPv4 App work with IPv6](#)

[Alternative Options for Prefix Length Handling](#)

# Benefits of this Library

---

The library was intended to satisfy the following primary goals:

- **Parsing of all host name and ipv4/ipv6 address formats in common usage** plus some additional formats (see below or see [javadoc for the IPAddressString class](#) for the extensive list)
- **Parsing and representation of subnets**, either those specified by network prefix length or those specified with ranges of segment values. For example, all strings in the list below represent the same IPv4 subnet:
  - with CIDR network prefix length: `1.2.0.0/16`
  - with mask: `1.2.0.0/255.255.0.0`
  - wildcard segments: `1.2.*.*`
  - range segments: `1.2.0-255.0-255`
  - range using inet\_aton format: `0x1.0x2.0x0-0xffff`
  - SQL-style single wildcards to end segments: `1.2.____.____`
  - IPv4 mapped IPv6: `::ffff:1.2.0.0/112`
  - Hexadecimal values: `0x01020000-0x0102ffff`
  - Octal values: `000100400000-000100577777`
- **Allow the separation of address parsing from host parsing.** In some cases you may have an address, in others you may have a host name, in some cases either one, so this supports all three options (for instance, when validating invalid input "1.2.3.a" as an address only, it will not be treated as a host with DNS lookup attempted in the way that `InetAddress.getByName` does)
- **Allow control over which formats are allowed when parsing**, whether IPv4/6, or subnets, or inet\_aton formats, and so on.
- **Produce all common address strings of different formats for a given IPv4 or IPv6 address and produce collections of such strings**

Some addresses can have hundreds of thousands of potential string representations (when you consider hex capitalization, ipv6 compression, and leading zeros, the various IPv4 and IPv6 formats, and combinations of all the above), although there are generally a handful of commonly used formats. Generating these strings, can help when searching or matching addresses in databases or text.
- **Support parsing of all common MAC Address** formats in usage and produce all common MAC address strings of different formats
- **Integration of MAC Address with IPv6** with standard conversions
- **Integration of IPv4 Address with IPv6** through common address conversions
- **Polymorphism** is a key goal. The library maintains an address framework of interfaces and abstract classes that allow most library functionality to be independent of address type of version, whether IPv4, IPv6 or MAC. This allows for code which supports both IPv4 and IPv6 transparently. You can write generic non-version-specific code to validate addresses, connect to addresses, produce address strings, mask addresses, etc. You can make use of the address framework which is agnostic towards address version.

- **Thread-safety and immutability.** The core types (host names, address strings, addresses, address sections, address segments, address ranges) are all immutable (like `java.lang.String` or `java.lang.Integer` instances). They do not change their underlying value. For multi-threaded apps this is valuable.
- **Address modifications**, such as altering prefix lengths, masking, splitting into sections and segments, splitting into network and host sections, reconstituting from sections and segments
- **Address operations**, such as obtaining the prefix block subnet for a prefixed address, iterating, spliterating or streaming through subnets, iterating, spliterating or streaming through prefix blocks, incrementing and decrementing addresses by integer values, reversing address bits for endianness or DNS lookup, set-subtracting subnets from other subnets, subnetting, intersections of subnets, merging subnets, checking containment of addresses in subnets, listing subnets covering a span of addresses
- **Sorting and comparison** of host names, addresses, address strings and subnets
- **Address data structures** including the trie, associative trie, corresponding sets and maps, providing additional address operations such as containment operations, sorting operations, prefix block operations, and subnet traversals
- **Integrate with the java primitive types and the standard java classes** `java.net.InetAddress` , `java.net.Inet6Address` , `java.net.Inet4Address` , and `java.math.BigInteger` .
- **Making address manipulations easy**, so you do not have to worry about longs/ints/shorts/bytes/bits, signed/unsigned, sign extension, ipv4/v6, masking, iterating, and other implementation details.

## Code Examples

---

This document provides in-depth and extensive documentation for the library, and includes some code snippets throughout. However, for common use-cases, you may wish to go straight to the wiki [code examples which cover a wide breadth of common use-cases](#). The code examples are focused more on common use-cases and operations, while this document is focused more on covering all areas in more detail with more words. This document focuses more on one area at a time, while the examples are useful in showing how to combine the functionality to achieve various end results.

## Supported IP Address Parsing Formats

---

This includes, those supported by the well-known routines `inet_aton` and `inet_pton`, the subnet formats listed above, all combinations of the above, and others:

- all the formats supported by `inet_pton` and `inet_aton`
- all the formats supported by `nmap`
- all the formats produced by `netstat` involving hosts/addresses with ports/services
- the subnet formats listed above, whether prefixed, masked, wildcards, ranges
- IPv6 canonical, compressed `1::1` , mixed `1:2:3:4:5:6:1.2.3.4` , `[bracketed]:port` , `::1:service` , and so on
- Hex values
- IPv6 base 85 values

- \* represents all addresses both ipv4 and ipv6
- /x where x is a positive integer, with no associated address, is interpreted as the network mask for prefix length x
- " " the empty string is considered the default loopback

For a more detailed list of formats parsed, some examples are below, or see the [javadoc for IPAddressString](#) .

## Subnet formats

- **CIDR (Classless Inter-Domain Routing) prefix length subnets**  
Adding the prefix length /x creates the address or subnet for that network prefix length, depending upon the configuration for prefix handling. The subnet 1.2.0.0/16 is the set of all addresses starting with 1.2
- **Wildcard \* \_ and range – subnets:**
  - \* denotes all possible values in one or more segments, so 1.\*.\* or just 1.\* is equivalent to 1.0.0.0/8
  - 0-1 denotes the range from 0 to 1
  - \_ replaces any digit at the end of a segment, for example 1\_ represents 10 to 19 in decimal or 10 to 1f in hex
- **Combinations:** Applying a prefix length to a subnet simply applies the prefix to every element of the subnet. 1\*.0.0/16 is the same subnet of addresses as 1.0.0.0/8

For a more detailed list of formats parsed, some examples are below, or see the [javadoc for IPAddressString](#) .

## Core classes

The core classes are **HostName**, **IPAddressString**, and **MACAddressString** along with the **Address** base class and its subclasses **IPAddress**, **IPv4Address**, **IPv6Address**, and **MACAddress**, as well as the sequential address class **IPAddressSeqRange** and its subclasses **IPv4AddressSeqRange** and **IPv6AddressSeqRange**. If you have a textual representation of an IP address, then start with **HostName** or **IPAddressString** . If you have numeric bytes or integers, then start with **IPv4Address** , **IPv6Address** or **MACAddress** . Note that address instances can represent either a single address or a subnet. If you have either an address or host name, or you have something with a port or service name, then use **HostName**.

## Parse String Representation of IP Address or Host Name

**IPAddressString** is used to convert. You can use one of **getAddress** or **toAddress** , the difference being whether parsing errors are handled by exception or not.

```
IPAddress address = new IPAddressString("1.2.3.4").getAddress();
if(address != null) {
    //use address here
}
```

or

```
String str = "1.2.3.4";
try {
```

```

    IPAddress address = new IPAddressString(str).toAddress();
    //use address here
} catch (AddressStringException e) {
    String msg = e.getMessage();//detailed message indicating issue
}

```

If you have either a host name or an address, you can use `HostName`:

```

public static void main(String[] args) {
    check(new HostName("::1"));
    check(new HostName("*"));
    check(new HostName("a.b.com"));
}

static void check(HostName host) {
    if(host.isAddress()) {
        System.out.println("address: " +
            host.asAddress().toCanonicalString());
    } else if(host.isAddressString()) {
        System.out.println("address string with ambiguous address: " +
            host.asAddressString());
    } else {
        System.out.println("host name with labels: " +
            Arrays.asList(host.getNormalizedLabels()));
    }
}

```

Output:

```

address: ::1
address string with ambiguous address: *
host name with labels: [a, b, com]

```

## Format Examples

Many formats are supported. For instance, the address `1:2:3:0:0:6::` can be represented many ways as shown.

```

static void parse(String formats[]) {
    for(String format : formats) {
        System.out.println(new IPAddressString(format).getAddress());
    }
}

static void parseHost(String formats[]) {
    for(String format : formats) {
        System.out.println(new HostName(format).getAddress());
    }
}

public static void main(String[] args) {
    String formats[] = {
        "1:2:3:0:0:6::",
        "1:2:3:0:0:6:0:0",
        "1:2:3::6:0:0",
        "0001:0002:0003:0000:0000:0006:0000:0000",
        "1:2:3::6:0.0.0.0",
        "1:2:3:0:0:6::",
    };
}

```

```

        "008JQW0V70(=61h*;$LC",
        "0x00010002000300000000000600000000"
    };
    parse(formats);

    String hostFormats[] = {
        "[1:2:3:0:0:6::]",
        "[1:2:3:0:0:6:0:0]",
        "[1:2:3::6:0:0]",
        "[0001:0002:0003:0000:0000:0006:0000:0000]",
        "[1:2:3::6:0.0.0.0]",
        "[1:2:3:0:0:6::]",
        "[008JQW0V70(=61h*;$LC)",
        "[0x00010002000300000000000600000000]",
        "0.0.0.0.0.0.0.6.0.0.0.0.0.0.0.0.3.0.0.0.2.0.0.0.1.0.0.0.ip6.arpa",
        "1-2-3-0-0-6-0-0.ipv6-literal.net"
    }
    parseHost(hostFormats);
}

```

Output:

```

1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0
1:2:3::6:0:0

```

Subnet strings are supported as well, using CIDR prefix notation or characters indicating range ('-' for a specific range or '\*' for full-range segments).

For instance, the subnet ffff::/104 can be represented many ways:

```

static void parse(String formats[]) {
    for(String format : formats) {
        System.out.println(new
            IPAddressString(format).getAddress().assignPrefixForSingleBlock());
    }
}

static void parseHost(String formats[]) {
    for(String format : formats) {
        System.out.println(new
            HostName(format).getAddress().assignPrefixForSingleBlock());
    }
}

public static void main(String[] args) {
    String prefixedFormats[] = {
        "ffff::/104",
        "ffff:0:0:0:0:0:0:0/104",
        "ffff:0000:0000:0000:0000:0000:0000:0000/104",
    }
}

```



## Delimited Segments

The subnet formats allow you to specify ranges of values. However, if you wish to parse addresses in which values are delimited, then you can use the methods `parseDelimitedSegments(String)` and `countDelimitedAddresses(String)` in `IPAddressString`. The method `parseDelimitedSegments` will provide an iterator to traverse through the individual addresses.

For example, given `"1,2.3.4,5.6"` `parseDelimitedSegments` will iterate through `"1.3.4.6"`, `"1.3.5.6"`, `"2.3.4.6"` and `"2.3.5.6"`. You can then construct `IPAddressString` instances from those individual strings for parsing.

## Address or Host Name Validation Options

Validation options allow you to restrict the allowed formats, whether you wish to support just IPv4 or IPv6, or whether you wish to support just single addresses, or whether you wish to allow different variants.

For IP addresses, there is the class `IPAddressStringParameters` which can be passed into `IPAddressString`, and for host names there is the class `HostNameParameters` which can be passed into `HostName`.

You constrict validation options using the appropriate builder classes that are nested classes in the options classes. Here is an example constructing host options along with nested address options within:

```
HostNameParameters HOST_OPTIONS_EXAMPLE = new HostNameParameters.Builder().
    allowEmpty().
    setEmptyAsLoopback(false).
    setNormalizeToLowercase(true).
    allowBracketedIPv6(true).
    allowBracketedIPv4(true).
    getAddressOptionsBuilder().
        allowPrefix(true).
        allowMask(true).
        setRangeOptions(RangeParameters.WILDCARD_AND_RANGE).
        allow_inet_aton(true).
        allowEmpty(false).
        setEmptyAsLoopback(false).
        allowAll(false).
        allowPrefixOnly(false).
    getIPv4AddressParametersBuilder().
        allowPrefixLengthLeadingZeros(true).
        allowPrefixesBeyondAddressSize(false).
        allowWildcardedSeparator(true).
    getParentBuilder().
    getParentBuilder().
    toParams();
```

The default options used by the library are permissive and not restrictive.

## Host Name or Address with Port or Service Name

For an address or host with port or service name, use `HostName`. IPv6 addresses with ports should appear as `[ipv6Address]:port` to resolve the ambiguity of the colon separator, consistent with RFC 2732, 3986, 4038 and other RFCs. However, this library will parse IPv6 addresses without the brackets. You can use the `"expectPort"` setting of `HostNameParameters` to resolve ambiguities when the brackets are absent.



```

HostName hostName = new HostName("::1:80");
System.out.println("host: " + hostName.getHost() + " address: " +
    hostName.getAddress() + " port: " + hostName.getPort());

hostName = new HostName("localhost:80");
System.out.println("host: " + hostName.getHost() + " port: " +
    hostName.getPort());

hostName = new HostName("127.0.0.1:80");
System.out.println("host: " + hostName.getHost() + " address: "
    + hostName.getAddress() + " port: " + hostName.getPort());

```

Output:

```

host: ::1 address: ::1 port: 80
host: localhost port: 80
host: 127.0.0.1 address: 127.0.0.1 port: 80

```

## IP Version Determination and IPv4/v6 Conversion

With an `IPAddress` or `IPAddressString` object, you can check the version with `isIPv4()` and `isIPv6()` and get the appropriate subclass with `toIPv4()` or `toIPv6()`. You can also make use of `isIPv4Convertible()` and `isIPv6Convertible()` to do further conversions if the address is IPv4-mapped, or you can use your own instance of `IPAddressConverter` for some other suitable conversion (IPv4-translated, 6to4, 6over4, IPv4 compatible, or other).

```

IPAddressString str = new IPAddressString("::ffff:1.2.3.4");
if(str.isIPv6()) {
    IPv6Address ipv6Address = str.getAddress().toIPv6();
    System.out.println(ipv6Address.toMixedString());
    if(ipv6Address.isIPv4Convertible()) {
        IPv4Address ipv4Address = ipv6Address.toIPv4();
        System.out.println(ipv4Address.toNormalizedString());
    }
}

```

Output:

```

::ffff:1.2.3.4
1.2.3.4

```

Should you wish to change the default IPv4/IPv6 conversions from IPv4 mapped to something else, you can override the pair of methods `toIPv4()` and `isIPv4Convertible()` in your own `IPv6Address` subclass and/or the pair of methods `toIPv6()` and `isIPv6Convertible()` in your own `IPv4Address` subclass.

## Prefixed Addresses and toHostAddress()

This library will parse CIDR prefix IP addresses such as `10.1.2.3/24`. On its own, that string can be interpreted both as an individual address, or as a prefix block with prefix `10.1.2`. With other libraries, the ambiguity is resolved by the method, function, or class used for parsing. However, this library uses the same types and methods for both subnets and individual addresses, which can result in much cleaner and more flexible code.

Network addresses (addresses with host that is zero) like 10.1.2.0/24 and a:b:c:d::/64 are parsed as the prefix block of addresses for the indicated prefix length. If the host is not zero, like 10.1.2.3/24, then it is parsed as an individual address, an address with an associated prefix length. This applies to both IPv4 and IPv6.

The same rule applies to strings in which the address is a subnet in which the subnet lower and upper values have zero hosts, like 10.1.2-3.0/24 or 10.1.2.2-6/31. Both of those examples will become CIDR prefix blocks when parsed, the first with 512 addresses, and the second with 10 addresses.

If you have an individual address or a subnet with prefix length and you want to get the address instance representing the entire block for that prefix, you can call `toPrefixBlock()`. In the reverse direction, given an entire block for a prefix, you can get the IPv4 network address or the IPv6 anycast address by calling `getLower()`, or you can use `getHostAddress()` or `toHostAddress()` when parsing, to get the address without the prefix.

With `getHostAddress()` or `toHostAddress()` you will get the host address 10.1.2.3 when parsing 10.1.2.3/24. You can also get the prefix length by calling `getNetworkPrefixLength()` on the `IPAddressString` or `IPAddress` instance.

The example code below shows some of these methods in use. Take note of the count for each string parsed, indicating whether you have a subnet or address.

```
static void printPrefixedAddresses(String addressStr) {
    IPAddressString ipAddressString = new IPAddressString(addressStr);
    IPAddress address = ipAddressString.getAddress();
    System.out.println("count: " + address.getCount());
    IPAddress hostAddress = ipAddressString.getHostAddress();
    IPAddress prefixBlock = address.toPrefixBlock();
    Integer prefixLength = ipAddressString.getNetworkPrefixLength();
    System.out.println(address);
    System.out.println(address.toCanonicalWildcardString());
    System.out.println(hostAddress);
    System.out.println(prefixLength);
    System.out.println(prefixBlock);
    System.out.println();
}

printPrefixedAddresses("10.1.2.3/24"); // individual address
printPrefixedAddresses("10.1.2.0/24"); // network
```

Output:

```
count: 1
10.1.2.3/24
10.1.2.3
10.1.2.3
24
10.1.2.0/24
```

```
count: 256
10.1.2.0/24
10.1.2.*
10.1.2.0
24
10.1.2.0/24
```

Typically, the segments or other punctuation identify a string as a host name, as an IPv4 address, or as an IPv6 address. The parser will also parse single segment values or a range of single segment values.

With non-segmented addresses, ambiguity between IPv4 and IPv6 is resolved by the number of digits in the string. The number of digits is 32 for IPv6 hexadecimal, 20 for IPv6 base 85 (see RFC 1924), and 11 or less for IPv4, which can be octal, hexadecimal, or decimal. For IPv4, digits are presumed decimal unless preceded by 0x for hexadecimal or 0 for octal, as is consistent with the `inet_aton` routine. For IPv6, 32 digits are considered hexadecimal and the preceding 0x is optional.

```
IPAddressString ipAddressString = new
IPAddressString("4)+k&C\#VzJ4br>0wv%Yp"); // base 85
IPAddress address = ipAddressString.getAddress();
System.out.println(address);

ipAddressString = new
IPAddressString("1080000000000000000080800200c417a"); // hex IPv6
address = ipAddressString.getAddress();
System.out.println(address);

ipAddressString = new IPAddressString("0x01020304"); // hex IPv4
address = ipAddressString.getAddress();
System.out.println(address);

ipAddressString = new IPAddressString("000100401404"); // octal IPv4
address = ipAddressString.getAddress();
System.out.println(address);
```

Output:

```
1080::8:800:200c:417a
1080::8:800:200c:417a
1.2.3.4
1.2.3.4
```

When parsing a range of single-segment values, it might not be possible to represent the range as a series of segments of range values, which is what is needed to be represented by an `IPv6Address` of 8 segment ranges, or an `IPv4Address` of 4 segment ranges. However, the string can still be parsed, and the parsed result can be obtained using `toDivisionGrouping`, `getDivisionGrouping`, `getSequentialRange`, or `toSequentialRange`. The former two methods provide an exact representation of the string divisions, and the latter provides an `IPAddressSeqRange` instance with the range of addresses from the lower to the upper value of the range expressed by the string, from which a series of `IPAddress` instances can be obtained using `spanWithPrefixBlocks` or `spanWithSequentialBlocks`.

## Parse Special Host Names – Reverse DNS Host Name, IPv6 Literal UNC Host Name

A couple of standardized host formats are recognized, namely the reverse DNS host format, and the UNC IPv6 literal host format.

```

HostName hostName = new
HostName("a.7.1.4.c.0.0.2.0.0.8.0.8.0.0.0.0.0.0.0.0.0.0.0.0.0.8.0.1.ip6.arpa");
System.out.println(hostName.asAddress());

hostName = new HostName("4.3.2.1.in-addr.arpa");
System.out.println(hostName.asAddress());

hostName = new

```

```
HostName("1080-0-0-0-8-800-200c-417a.ipv6-literal.net");
System.out.println(hostName.asAddress());
```

Output:

```
1080::8:800:200c:417a
1.2.3.4
1080::8:800:200c:417a
```

A couple of methods in `HostName` are available as well:

```
public boolean isUNCIPv6Literal()
public boolean isReverseDNS()
```

## Parse IPv6 Zone or Scope ID

The IPv6 zone or scope ID is recognized, denoted by the `'%'` character. It can be retrieved by `IPv6Address.getZone()` .

```
IPAddress addr = new IPAddressString("::%eth0").getAddress();
if(addr.isIPv6()) {
    System.out.println(addr.toIPv6().getZone());
}
```

Output:

```
eth0
```

## IP Address and Numeric Values

---

In addition to the range of string formats that can be parsed to produce `IPAddress` instances, you can also obtain instances of `IPAddress` from byte arrays, `java.net.InetAddress` , arrays of address segments, or individual integer segment values.

For IPv4 you have the additional option of constructing an address from a 32-bit integer. For IPv6, you have the additional options of constructing from a `java.math.BigInteger` and constructing from MAC addresses.

Once you have an `IPAddress` instance, there are methods to convert to bytes, to a `java.math.BigInteger` , to sections, to subnets or network prefixes, to masks, to `java.net.InetAddress` , to different string representations, and so on.

When constructing IP addresses or sections, the same rules regarding zero-hosts applies as when parsing from strings. Network addresses (addresses with host that is zero) like `10.1.2.0/24` and `a:b:c:d::/64` are constructed as the prefix block of addresses for the indicated prefix length. If the host is not zero, like `10.1.2.3/24`, then it is constructed as an individual address, an address with an associated prefix length. This applies to both IPv4 and IPv6.

The same rule applies to subnets where the lower and upper values have zero hosts, like `10.1.2-3.0/24` or `10.1.2.2-6/31`. Both of those examples will become CIDR prefix blocks when constructed.

Should you wish to get the individual address or section with a zero host, you can construct without the prefix length and then apply the prefix length afterwards, or you can use `getLower()` after construction.

## Networks

---

Each of the different address types (IPv6, IPv4, MAC) has an associated singleton network object. The network objects are used for caching, for configuration, and have methods for obtaining masks and loopbacks. Each of the network objects have singleton creator objects that are used to create addresses, sections, and segments, and those creator objects perform caching of these address components for efficient memory usage and performance. All internally created address components and addresses are created by the creator object, whether by the string parsing engine or by address object manipulation.

The `defaultIPv6Network()`, `defaultIPv4Network()`, and `defaultMACNetwork()` methods in `Address` provide access to the respective network objects. Each network's associated creator object is accessible from `getAddressCreator()`.

### Using your own address classes

Users can create their own subclasses of the address and address component classes (eg your own subclasses of `IPv6Address`, `IPv6AddressSection` and/or `IPv6AddressSegment`).

When doing so, you may wish to use your own network classes as well. To do so, you

- override `getNetwork()`
- override `getIPv6Network()`, `getIPv4Network()`, and `getMACNetwork()`, if any of those methods exist in the class you are subclassing. Several of those methods may exist in the same class because address components sometimes create components of other address versions or types (eg embedding IPv4 in IPv6, or MAC integration with IPv6).

Once you have done so, your new address classes are using your own network classes as well.

When using your own network classes, you can override the method `createAddressCreator()` or the method `getAddressCreator()` in `AddressNetwork`, and in so doing provide your own creator objects, in which you can override any or all of the creation methods to use your own classes.

You can make the string parsing engine use your own address or network classes. To do this, your code should create your own string validation objects using the builder classes `MACAddressStringParameters.Builder`, `IPAddressStringParameters.Builder` and `HostNameParameters.Builder`. With `HostNameParameters.Builder` you must get the `IPAddressStringParameters.Builder` sub-builder by calling `getAddressOptionsBuilder()`. With `IPAddressStringParameters.Builder` you must get the IPv4 or IPv6 sub-builder by calling one of `getIPv6Parameters()` or `getIPv4Parameters()`. You then call `setNetwork()` on the builder or sub-builder objects to supply your own network. The initial builder can create your own validation objects with `toParams()`. Those validation objects can be supplied to the appropriate constructor of `IPAddressString` or `MACAddressString`, and the subsequent string parsing with those instances will use the supplied networks.

## Prefix Length Handling

---

Prefix lengths are supported both in parsing, as indicated in the section above on parsing, as well as when directly constructing addresses or sections. Addresses and sections store their prefix lengths and the prefix is incorporated in

numerous address operations as well as when producing strings. For instance, an address will provide the network section upon calls to `getNetworkSection()` in `IPAddress`, and will supply the prefix length when calling `getNetworkPrefixLength()`.

Given an address with no prefix length, you can convert to an address with prefix length using the methods `assignPrefixForSingleBlock()` or `assignMinPrefixForBlock()` in `IPAddress`, or any of the methods that allow you to set any given prefix length directly such as `setPrefixLength(int)`.

## Prefix Length and Equality

In this library, the subnet with prefix length `10.1.2.*`/`24` is equivalent the non-prefixed address `10.1.2.*` as they both contain the same set of addresses. In other words, when it comes to equality or comparison, the prefix length has no effect.

## Address Sections

---

Addresses can be broken up into sections, and reconstituted from sections. A section is a series of segments. You can get the section for the full address by calling `getSection()`, or you can get subsections by calling one of the variants, either `getSection(int)` or `getSection(int, int)`, which return a subsection spanning the given indices. You can also get the segments in an address by calling `getSegment(int)` or one of the variants of `getSegments`, either on the address or on a section of the address.

You can also reconstitute an address from a section or array of segments using the appropriate address constructor, if your section or array of segments has the correct number of segments for the address type.

### Host and Network Sections of IP Address

Use `getHostSection()` and `getNetworkSection()` to get the host and network sections of an IP address as indicated by prefix length.

```
IPAddress address = new IPAddressString("1.2.3.4").getAddress();
IPAddressSection network = address.getNetworkSection(16, true);
IPAddressSection host = address.getHostSection(16);
System.out.println(network.toCanonicalString());
System.out.println(host.toCanonicalString());
```

Output:

```
1.2/16
3.4
```

Once you have a section of an address, most of the same methods are available as those available with addresses themselves.

## IP Address Ranges

---

An `IPAddress` or `IPAddressString` instance can represent any individual address or any range of addresses in which each segment specifies its own range of sequential values. For `IPAddress`, the ranges are specified with the canonical number of segments for each address version or type, which is 4 segments for IPv4, 8 for IPv6, and either 6 or 8 for MAC.

Any prefix block of addresses can be specified this way (a block of addresses that includes all the hosts for a given prefix).

An `IPAddressString` can represent any such address string, even those that do not use the canonical number of segments. However, for ranges involving a non-canonical number of segments, converting to an `IPAddress` instance is not always possible (for example the IPv4 subnet 1.2-3.0-1000 cannot be expressed with 4 segments).

Not all `IPAddress` subnets are sequential. For instance, 1-2.3.4-5.6 is not sequential, since the address 1.3.4.6 is followed in the block by 1.3.5.6 while the next sequential address 1.3.4.7 is not part of the block. A sequential block is one in which the range is sequential. For a block to be sequential, the first segment with a range of values must be followed only by segments that cover all values. For instance, 1.2.3-4.\* is a sequential block, as well as 1:a-f:\*:\*:\*:\*, also writeable as 1:a-f:\*. Any prefix block is sequential and can also be expressed without the prefix length. For instance, the prefix block 1:2:3:4::/64 can be written without the prefix length as 1:2:3:4.\*.

You can convert a non-sequential block to a collection of sequential blocks using the `sequentialBlockIterator` method of `IPAddress`. If you wish to get the count of sequential blocks, use the method `getSequentialBlockCount`.

```
convertNonSequentialBlock("a:b:c:d:1:2-4:3-5:4-6");

static void convertNonSequentialBlock(String string) {
    IPAddressString addrString = new IPAddressString(string);
    IPAddress addr = addrString.getAddress();
    System.out.println("Initial range block is " + addr);
    BigInteger sequentialCount = addr.getSequentialBlockCount();
    System.out.println("Sequential range block count is " + sequentialCount);

    Iterator<? extends IPAddress> iterator = addr.sequentialBlockIterator();
    while(iterator.hasNext()) {
        System.out.println(iterator.next());
    }
}
```

Output:

```
Initial range block is a:b:c:d:1:2-4:3-5:4-6
Sequential range block count is 9
a:b:c:d:1:2:3:4-6
a:b:c:d:1:2:4:4-6
a:b:c:d:1:2:5:4-6
a:b:c:d:1:3:3:4-6
a:b:c:d:1:3:4:4-6
a:b:c:d:1:3:5:4-6
a:b:c:d:1:4:3:4-6
a:b:c:d:1:4:4:4-6
a:b:c:d:1:4:5:4-6
```

Not all sequential address ranges can be described by an instance of `IPAddress` or `IPAddressString`. One such example is the range of two IPv4 addresses from 1.2.3.255 to 1.2.4.0. One option is to represent the address range with just a single large segment covering the section of the address that has a range of values, such as in an instance of `IPAddressLargeDivisionGrouping`.

The better option is to use an `IPAddressSeqRange` instance, which provides a more general representation of address ranges and their associated operations. You can represent any sequential range of addresses with an `IPAddressSeqRange`.



instance.

Both `IPAddress` and `IPAddressSeqRange` implement the `IPAddrRange` interface. `IPAddressSeqRange` instances cover all ranges of addresses, while `IPAddress` instances cover all ranges of segments, including CIDR prefix blocks.

An `IPAddressSeqRange` instance can be converted to the minimal list of sequential blocks or prefix blocks that cover the same range of addresses, so that is a couple of ways to go from `IPAddressSeqRange` to `IPAddress` instances. To go in the reverse direction from `IPAddress` to `IPAddressSeqRange`, you can convert any `IPAddress` sequential block or prefix block to an `IPAddrRange` with the method `toSequentialRange`, and then you can join any collection of `IPAddrRange` instances with the join method in `IPAddressSeqRange`.

Here we show a code example of creating a sequential range, breaking it up into either prefix blocks or sequential blocks, and then merging those blocks to get the original sequential range.

```
spanAndMerge("2:3:ffff:5::", "2:4:1:5::");

static void spanAndMerge(String address1, String address2) {
    IPAddressString string1 = new IPAddressString(address1);
    IPAddressString string2 = new IPAddressString(address2);
    IPAddress addr1 = string1.getAddress();
    IPAddress addr2 = string2.getAddress();
    IPAddressSeqRange range = addr1.toSequentialRange(addr2);
    System.out.println("Original range of size " +
        range.getCount() + ": " + range);

    IPAddress result[] = range.spanWithPrefixBlocks();
    System.out.println("Prefix blocks: " + Arrays.asList*(result));

    IPAddress result2[] = range.spanWithSequentialBlocks();
    System.out.println("Sequential blocks: " +

    Arrays.asList(result2));
    List<IPAddressSeqRange> rangeList = new ArrayList<>();
    for(IPAddress a : result) {
        IPAddressSeqRange r = a.toSequentialRange();
        rangeList.add(r);
    }
    IPAddressSeqRange joined[] = IPAddressSeqRange.join(
        rangeList.toArray(new IPAddressSeqRange[rangeList.size()]));
    System.out.println("Merged prefix blocks back again: " +

    Arrays.asList(joined));
    rangeList.clear();
    for(IPAddress a : result2) {
        IPAddressSeqRange r = a.toSequentialRange();
        rangeList.add(r);
    }
    joined = IPAddressSeqRange.join(
        rangeList.toArray(new IPAddressSeqRange[rangeList.size()]));
    System.out.println("Merged sequential blocks back again: " +
        Arrays.asList(joined));
}
```

Output:

Original range of size 2417851639229258349412353: 2:3:ffff:5:: -> 2:4:1:5::

Prefix blocks: [2:3:ffff:5::/64, 2:3:ffff:6::/63, 2:3:ffff:8::/61, 2:3:ffff:10::/60, 2:3:ffff:20::/59, 2:3:1

Sequential blocks: [2:3:ffff:5-ffff:\*\*\*\*, 2:3-4:\*\*\*\*, 2:4:0:\*\*\*\*, 2:4:1:0-4:\*\*\*\*, 2:4:1:5:



Merged prefix blocks back again: [2:3:ffff:5:: -> 2:4:1:5::]

Merged sequential blocks back again: [2:3:ffff:5:: -> 2:4:1:5::]

As you can see in the example above, you can generally describe a range with fewer sequential blocks than prefix blocks.

## Address Tries

The trie data structure is particularly useful when working with addresses. For that reason this library includes compact binary address tries (aka compact binary prefix tree or binary radix trie, amongst other names). Tries provide efficient retrieval operations, hence the name *trie*, but what makes them additionally useful for addresses is the fact that prefix tries are organized based on the prefix bits of the keys, which are addresses in this case, mirroring the way that CIDR subnets and addresses are organized by prefix. So you can use tries for efficient subnet containment checks on many addresses or subnets at once in constant time (such as with a routing table), for efficient lookups, for efficiently dividing and subdividing subnets, for sorting addresses, and for traversing through subnets in different ways.

By also associating each trie node with a value, tries can be used for value lookups in which the keys are addresses.

Tries can also be used as the backing data structures for Java collections framework sets and maps.

When handling large numbers of addresses or CIDR prefix blocks, it can be much more efficient to use the trie data structure for common operations, the trie constructed in linear time proportional to the number of addresses, and then offering constant time containment and retrieval operations on all of the contained addresses or subnets at once.

An IPv6 trie constructed and then converted to a string:

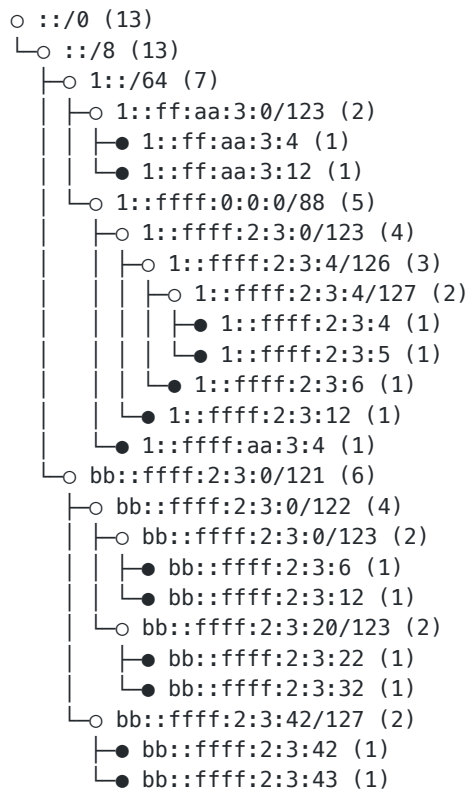
```
<T extends AddressTrie<A>, A extends Address> T populateTree(
    T trie, String addrStrs[], Function<String, A> creator) {
    for(String addrStr : addrStrs) {
        A addr = creator.apply(addrStr);
        trie.add(addr);
    }
    return trie;
}

String ipv6Addresses[] = {
    "1::ffff:2:3:5",
    "1::ffff:2:3:4",
    "1::ffff:2:3:6",
    "1::ffff:2:3:12",
    "1::ffff:aa:3:4",
    "1::ff:aa:3:4",
    "1::ff:aa:3:12",
    "bb::ffff:2:3:6",
    "bb::ffff:2:3:12",
    "bb::ffff:2:3:22",
    "bb::ffff:2:3:32",
    "bb::ffff:2:3:42",
    "bb::ffff:2:3:43",
};

IPv6AddressTrie ipv6Trie = populateTree(
    new IPv6AddressTrie(),
    ipv6Addresses,
    str -> new IPAddressString(str).getAddress().toIPv6());

System.out.println(ipv6Trie);
```


Output:




An IPv4 trie constructed and then converted to a string using the same polymorphic `populateTree` method:

```
String ipv4Addresses[] = {
    "1.2.3.4",
    "1.2.3.5",
    "1.2.3.6",
    "1.2.3.3",
    "1.2.3.255",
    "2.2.3.5",
    "2.2.3.128",
    "2.2.3.0/24",
    "2.2.4.0/24",
    "2.2.7.0/24",
    "2.2.4.3",
};
IPv4AddressTrie ipv4Trie = populateTree(
    new IPv4AddressTrie(),
    ipv4Addresses,
    str -> new IPAddressString(str).getAddress().toIPv4());
System.out.println(ipv4Trie);
```

Output:





A more compact non-binary representation of the same IPv4 trie:

```
System.out.println(ipv4Trie.toAddedNodesTreeString());
```

Output:

```
○ 0.0.0.0/0
├─● 1.2.3.3
├─● 1.2.3.4
├─● 1.2.3.5
├─● 1.2.3.6
├─● 1.2.3.255
├─● 2.2.3.0/24
│   └─● 2.2.3.5
│       └─● 2.2.3.128
├─● 2.2.4.0/24
│   └─● 2.2.4.3
└─● 2.2.7.0/24
```

## Partitioning Subnets

An address trie accepts individual addresses or CIDR prefix blocks subnets. While this covers most use-cases, there are `IPAddress` instances that cannot be added to a trie as-is. But any subnet can be subdivided or partitioned in various ways.

The `Partition` class encapsulates a partition of a subnet. It also provides a couple of methods that subdivide any subnet into individual addresses or prefix block subnets, which can then be inserted into a trie. Much like an iterator, a partition can be used once, so simply create another if that is necessary.

The two partition methods provided partition differently. `partitionWithSingleBlockSize` finds a maximal prefix block size and then iterates through a series of prefix blocks of that size. `partitionWithSpanningBlocks` uses any number of different prefix block sizes, which often results in a smaller number of blocks.

Here we partition an IPv4 subnet with `partitionWithSingleBlockSize` and then check for the partitioned elements in the trie:

```
String addr = "1.2.1-7.*";
IPv4AddressTrie trie = new IPv4AddressTrie();
IPv4Address subnet = new IPAddressString(addr).getAddress().toIPv4();
Partition.partitionWithSingleBlockSize(subnet).predicateForEach(trie::add);
boolean foundThemAll = Partition.partitionWithSingleBlockSize(subnet).predicateForEach(trie::contains);
```

```
System.out.println("all inserted: " + foundThemAll);
System.out.println(trie);
```

Output:

```
all inserted: true
```

```
○ 0.0.0.0/0 (7)
└○ 1.2.0.0/21 (7)
   ├──○ 1.2.0.0/22 (3)
   │   ├──● 1.2.1.0/24 (1)
   │   └──○ 1.2.2.0/23 (2)
   │       ├──● 1.2.2.0/24 (1)
   │       └──● 1.2.3.0/24 (1)
   └──○ 1.2.4.0/22 (4)
       ├──○ 1.2.4.0/23 (2)
       │   ├──● 1.2.4.0/24 (1)
       │   └──● 1.2.5.0/24 (1)
       └──○ 1.2.6.0/23 (2)
           ├──● 1.2.6.0/24 (1)
           └──● 1.2.7.0/24 (1)
```

Here we try the other partition method `partitionWithSpanningBlocks``:

```
trie = new IPv4AddressTrie();
Partition.partitionWithSpanningBlocks(subnet).predicateForEach(trie::add);
foundThemAll = Partition.partitionWithSpanningBlocks(subnet).predicateForEach(trie::contains);
System.out.println("all inserted: " + foundThemAll);
System.out.println(trie);
```

Output:

```
all inserted: true
```

```
○ 0.0.0.0/0 (3)
└○ 1.2.0.0/21 (3)
   ├──○ 1.2.0.0/22 (2)
   │   ├──● 1.2.1.0/24 (1)
   │   └──● 1.2.2.0/23 (1)
   └──● 1.2.4.0/22 (1)
```

The two tries illustrate how the two partitions differ.

## IP Address Operations

---

There are various methods for masking, obtaining subnets, and so on.

### Summary of IP Address Operations

Here is a summary, a non-exhaustive list, in no specific order, of operations for transforming addresses and subnets. Many of these methods are available not just for address instances, but also for sections, for sequential ranges, and a few are available for segments as well:

- **toPrefixBlock**: Provides the subnet for the entire prefix block using the existing or a supplied prefix length. You can use `isPrefixBlock` to determine if the operation would have no effect on the address because it is already a prefix block.
- **assignPrefixForSingleBlock** (formerly `toPrefixedEquivalent`): Converts a subnet address with no prefix length to an equivalent address with a prefix length. The resulting address will span the same range of values as the original. You can use `getPrefixLengthForSingleBlock` (formerly `getEquivalentPrefix`) to get the resultant prefix prior to the operation, or `isSinglePrefixBlock` (formerly `isRangeEquivalentToPrefix`) to determine if the operation would have no effect on the address because it is already a single prefix block.
- **assignMinPrefixForBlock** (formerly `toMinPrefixedEquivalent`): Converts an address to an equivalent address with the smallest possible prefix length. The resulting address will span the same range of values as the original. You can use `getMinPrefixLengthForBlock` (formerly `getMinPrefix`) to get the resultant prefix prior to the operation.
- **mergeToPrefixBlocks, mergeToSequentialBlocks**: Given a list of addresses or subnets, merges them into the minimal list of prefix block subnets or sequential block subnets
- **join**: Given a list of address ranges, merges them into the minimal list of address ranges
- **spanWithPrefixBlocks, spanWithSequentialBlocks**: Given a pair of addresses or subnets, finds the minimal list of prefix block subnets or sequential block subnets that span all addresses between the pair
- **append, prepend, replace**: Add or replace segments in sections and addresses. On addresses you must always maintain the correct number of segments, so only the replace operation is provided. You can obtain a section by either constructing one directly or by calling a variant of `getSection` on an address.
- **adjustPrefixBySegment**: Will create a larger network or smaller subnet address by increasing or decreasing the prefix length to the next segment boundary. If prefix length is increased, the additional prefix bits can be either zero or the full range of potential values.
- **withoutPrefixLength**: Convert to the same address but with no prefix length
- **removePrefixLength, adjustPrefixLength, setPrefixLength**: Add, remove or adjust prefix lengths by the indicated values. There are variants with a boolean argument to control whether extended prefix lengths have zeros for the added bits, or whether the bits retain their values from when they were on the other side of the prefix boundary. The default behavior when extending or shortening a prefix length is to insert zeros in the bits added to the prefix or removed from the prefix. `removePrefixLength` is like `withoutPrefixLength` but will convert the former host bits to zero.

When reconstituting the result, if the initial host bits have been changed to zero or remain zero, while the remaining host bits are zero or cover the full range, then the result will be a prefix block, just like when parsing strings with zero hosts or creating addresses or sections with zero hosts.

- **mask, maskNetwork, bitwiseOr, bitwiseOrNetwork**: apply masks to subnets and addresses. The "network" variants allow you to mask just the network portion and apply a prefix length, while the others remove any existing prefix length.
- **reverseBits, reverseBytes, reverseBytesPerSegment, reverseSegments**: useful for handling endianness (network byte order sometimes requires bytes be reversed), or DNS lookup, or other reasons for reversing bits and bytes.
- **subtract**: Computes the subnet difference, the set of addresses in the subnet but not in the argument subnet. Subtracts a given subnet from the receiver subnet, returning an array of subnets for the result.
- **intersect**: Computes the subnet conjunction, the set of addresses in both the receiver subnet and the argument subnet.
- **increment**: Given a positive or negative integer increment value, if the address is a subnet, provides the address that is the given increment into the list of subnet addresses. Any increment exceeding the count of addresses in the subnet is

simply added to the final address in the subnet. If the address is an individual address, adds the given increment to the address value to produce a new address.

- **toIPv4, toIPv6, toEUI:** Use either standard or customized address conversions to go from one address version/type to another
- **iterator:** iterates through the individual addresses of a subnet
- **prefixBlockIterator:** iterates through the individual prefix block subnets of a larger subnet

## Queries for Prefix Lengths and Prefix Blocks

In some cases you may need to know if you have a prefix block subnet, such as the /64 block a:b:c:d::/64 or the /16 block 1.2.0.0/16, in which the subnet contains all the addresses with the same prefix of a given length. The following methods allow you to query for prefix blocks and lengths of blocks.

- **containsPrefixBlock:** Answers whether the subnet contains the prefix block of addresses for a given prefix length (regardless of whether the address is assigned a different prefix length or has none assigned).
- **containsSinglePrefixBlock:** Like `containsPrefixBlock`, but queries if the subnet contains ~~is~~ the prefix block of addresses for a given prefix length. In other words, it answers whether the subnet contains the full block of address for just a single prefix
- **isPrefixBlock:** Like `containsPrefixBlock` but using the prefix length assigned to the address.
- **isSinglePrefixBlock:** Like `containsSinglePrefixBlock` but using the prefix length assigned to the address
- **getBlockMaskPrefixLength:** (formerly `getMaskPrefixLength`): Will return the prefix length for a network or host mask, if it is one
- **getPrefixLengthForSingleBlock:** Returns a prefix length for which the range of this segment grouping matches exactly the block of addresses for that prefix, if such a prefix length exists, and if it does, it will match the result of `getMinPrefixLengthForBlock`
- **getMinPrefixLengthForBlock:** Returns the smallest prefix length such that this address division series includes the block of addresses for that prefix, which is the bit count for individual addresses.

## Iterators, Spliterators and Streams

Various iterators, spliterators and streams are available for traversing through subnets in different ways.

- **iterator, spliterator, stream:** Traverses through all address items. Use `getCount` to get the count.
- **nonZeroHostIterator:** Like `iterator`, but if the subnet has a prefix length, then this skips all zero-valued hosts. This can be useful because the zero-valued host in a subnet typically represents the subnet as a whole and is typically not used as an individual address. Use `getNonZeroHostCount` to get the iterated count.
- **prefixBlockIterator, prefixBlockSpliterator, prefixBlockStream:** If the subnet has a prefix length, then this traverses through the prefixes, with each traversed item being the prefix block for each prefix. If no prefix length, same as `iterator / spliterator / stream`. Use `getPrefixCount` to get the count.
- **prefixIterator, prefixSpliterator, prefixStream:** If the subnet has a prefix length, then this iterates through the prefixes, with each iterated item including all the addresses in the original iterated subnet which have the same prefix. For all

traversed items except the lowest and highest this is the prefix block, while the lowest and highest prefix might have fewer addresses than the prefix block. If no prefix length, same as `iterator / spliterator / stream`. Use `getPrefixCount` to get the count.

- **blockIterator, blockSpliterator, blockStream:** Traverses through the segment values of the higher segments up until a given segment index. If the segment index is 0, produces just a single value, the original address. If the segment index is the segment count, same as `iterator`. Like `prefixIterator / prefixSpliterator / prefixStream`, for each segment prefix, includes all addresses in the original iterated subnet which have that prefix. Use `getBlockCount` to get the count.
- **sequentialBlockIterator, sequentialBlockSpliterator, sequentialBlockStream:** Traverses through the segment values of the higher segments up until a specific segment index. The segment index is chosen to be the largest segment index which allows for all iterated address blocks to be sequential. Use `getSequentialBlockCount` to get the iterated count.
- **segmentIterator, segmentsSpliterator, segmentsStream:** Traverses through all address items, similar to `iterator / spliterator / stream`, but using only segment arrays. Use `getCount` to get the count.

## Mask and Prefix Length Operations

The methods `mask`, `maskNetwork`, `bitwiseOr`, `bitwiseOrNetwork`, `removePrefixLength`, `adjustPrefixBySegment`, `adjustPrefixLength`, `applyPrefixLength`, and `setPrefixLength` allow you to apply and adjust prefix lengths, and to apply masks, to individual addresses/sections and subnets.

When applying an operation to a subnet, the operation is applied to every member of the subnet, so the result must be something representable with sequential segment ranges, otherwise you will get `IncompatibleAddressException` (formerly `AddressTypeException`). For instance, masking the subnet block of 255 addresses `0.0.0.0/24` with `0.0.0.128` results in the two addresses `0.0.0.0` and `0.0.0.128` which is not sequential. However, this will not happen when using standard masking and subnetting techniques typical with IPv4 and IPv6 routing, so in the usual cases no exception will be thrown.

The following code demonstrates how to get the lowest address in a prefixed subnet using any one of three methods: `getLowest`, `mask`, or `removePrefixLength`.

```
String addr = "1.2.3.4";
IPAddress address = new IPAddressString(addr).getAddress();
int prefixLength = 16;
IPAddress maskWithPreLen =
    new IPAddressString("/" + prefixLength).getAddress(address.getIPVersion());
IPAddress mask = address.getNetwork().getNetworkMask(16, false);

System.out.println("mask with prefix length " + maskWithPreLen);
System.out.println("mask " + mask);

IPAddress maskedAddress = address.mask(mask);

System.out.println("address " + address + " masked " + maskedAddress);

// create the subnet

IPAddress subnet =
    address.applyPrefixLength(prefixLength).toPrefixBlock();

// mask

IPAddress maskedSubnet = subnet.mask(mask);

System.out.println("subnet " + subnet + " masked " +
```

```

maskedSubnet);
System.out.println("equals: " + maskedAddress.equals(maskedSubnet));

// getLower

IPAddress lowestInSubnet = subnet.getLower();

System.out.println("lowest in subnet " + lowestAddressInSubnet);
System.out.println("lowest in subnet no prefix " + lowestInSubnet.removePrefixLength(false));
System.out.println("equals: " + maskedAddress.equals(lowestAddressInSubnet));

// removePrefixLength

IPAddress prefixRemoved = subnet.removePrefixLength(true);
System.out.println("prefix removed " + prefixRemoved);

```

Output:

```

mask with prefix length 255.255.0.0/16
mask 255.255.0.0
address 1.2.3.4 masked 1.2.0.0
subnet 1.2.0.0/16 masked 1.2.0.0
equals: true
lowest in subnet 1.2.0.0/16
lowest in subnet no prefix 1.2.0.0
equals: true
prefix removed 1.2.0.0

```

## Polymorphism

Simply change the string "1.2.3.4" in the code above to an IPv6 address like a:ffff:b:c:d::f and the code works all the same.

Output:

```

mask with prefix length ffff::/16
mask ffff::
address a:ffff:b:c:d::f masked a::
subnet a::/16 masked a::
equals: true
lowest in subnet a::/16
lowest in subnet no prefix a::
equals: true
prefix removed a::

```

## Subnetting

Subnetting can be accomplished using various address manipulation methods. Given a prefixed IP address, you can extend the prefix length and insert bits for an extended prefix and new subnet of the previous address block, as shown in the following example.

```

int originalPrefix = 18, adjustment = 4;
IPAddress address = new IPAddressString("207.0.64.0").
    getAddress();
IPAddress address2 = address.setPrefixLength(originalPrefix);

```



```

IPAddress prefixExtension = new
IPAddressString("0.0.4.0").getAddress();
IPAddress subnet1 = address2.adjustPrefixLength(adjustment); // extend the prefix length
System.out.println(subnet1);
IPAddress subnet2 =
    subnet1.bitwiseOrNetwork(prefixExtension,
        originalPrefix + adjustment); // adjust the extended prefix
System.out.println(subnet2);

```

Output:

```

207.0.64.0/18
207.0.64.0/22
207.0.68.0/22

```

Here is the same subnetting operation using segment replacement:

```

IPv4Address address = new IPAddressString("207.0.64.0/18").getAddress().toIPv4();
IPv4AddressSection replacementSection =
    new IPAddressString("0.0.68.0/22").getAddress().toIPv4().getSection(2);
IPAddress subnet = new IPv4Address(address.getSection().replace(2, replacementSection));
System.out.println(subnet);

```

Output:

```

207.0.68.0/22

```

Alternatively, you can use the `prefixBlockIterator()` method to get a list of subnets when adjusting the prefix:

```

IPAddress subnet = new
IPAddressString("192.168.0.0/28").getAddress();
IPAddress newSubnets = subnet.setPrefixLength(subnet.getPrefixLength() + 2, false);
System.out.println(newSubnets);

TreeSet<IPAddress> subnetSet = new TreeSet<IPAddress>();
Iterator<? extends IPAddress> iterator = newSubnets.prefixBlockIterator();
while (iterator.hasNext()) {
    subnetSet.add(iterator.next());
}
System.out.println(subnetSet);

```

Output:

```

192.168.0.0-12/30
[192.168.0.0/30, 192.168.0.4/30, 192.168.0.8/30, 192.168.0.12/30]

```

## Trie Operations

Most of these operations are methods that operate directly on `AddressTrie` or `AssociativeAddressTrie` instances, but some are available from their associated `AddressTrieSet` or `AddressTrieMap` instances available from `asSet` or `asMap`, in which case the methods would still be operating on the same backing trie.

- **add, addNode, addTrie, addAll, put, putNew, putNode, putTrie, putAll:** Insert one or more nodes into the trie for the given individual address(es) or CIDR prefix block(s). The put methods associate a value with the node.
- **get, getNode, getAddedNode, contains, containsAll, containsKey:** Retrieves or checks for the existence of one or more nodes in the trie whose keys are the given individual addresses or CIDR prefix blocks.
- **remove, removeAll:** Removes from the trie the nodes whose keys are the given individual addresses or CIDR prefix blocks, if any.
- **ceilingAddedNode, floorAddedNode, higherAddedNode, lowerAddedNode, ceiling, floor, higher, lower, ceilingKey, floorKey, higherKey, lowerKey, ceilingEntry, floorEntry, higherEntry, lowerEntry:** Finds nodes whose keys are closest to the given individual address or CIDR prefix block. These are implementations of methods defined by `java.util.NavigableSet` and `java.util.NavigableMap`.
- **firstNode, firstAddedNode, lastNode, lastAddedNode:** Retrieves the node for the lowest or highest valued key/address in the trie, according to the trie order. The “added” variants give the same answer in most cases, because the first and last node is usually an added node, the exceptions being when all the nodes are bigger than the non-added root, when all the nodes are smaller than the non-added root, or both conditions are true and the trie is empty.
- **elementContains, elementsContaining:** Retrieves or checks for the existence of a node in the trie whose key/subnet/address contains the given individual address or CIDR prefix block. Note the distinction with `contains`, which refers to containment in the data structure itself, rather than these methods referring to subnet containment by a key/subnet/address in the data structure.
- **elementsContainedBy, removeElementsContainedBy:** Retrieves or removes the sub-trie (if any) contained by the given individual address or CIDR prefix block. If the given argument is an individual address, then the returned or removed sub-trie can only be a single node or null.
- **headMap, tailMap, subMap, headSet, tailSet, subSet:** Creates a new `AddressTrieMap` or `AddressTrieSet` with a restricted range, created from the original `AddressTrieMap` or `AddressTrieSet`, with the new set or map backed by the same trie. Since it is backed by the same trie, changes to the new set or map will affect the original set or map and the shared trie, and vice-versa. After creating the new set or map, you can use `asTrie` on that new set or map to give you a new trie (and a new associated set or map), allowing you to make changes without affecting the original.
- **size, nodeSize:** returns the number of elements (added nodes) or the number of nodes (both added and auto-generated) in the trie. The `size` operation is constant time unless you are using a head/tail/sub `AddressTrieMap` or `AddressTrieSet` with a restricted range, in which case it is linear time like `nodeSize`.
- **remap, remapIfAbsent:** allows for changes to trie mappings based on the existing mappings, without having to do two lookups, one to examine the trie and a second to change the trie based on the results of the first lookup. Avoiding two lookups can be especially beneficial with tries being accessed by multiple threads, allowing for finer locking.
- **compute, computeIfAbsent, computeIfPresent, getOrDefault, merge, putIfAbsent:** map operations specified by `java.util.Map` and implemented in `AddressTrieMap` for efficiency and convenience, avoiding multiple lookups into the trie. Some of these use `remap` or `remapIfAbsent`.
- **constructAddedNodesTree:** Produces an associative trie from a trie. The produced trie maps the root and the added nodes to their added sub-nodes. This is essentially creating a non-binary tree structure from the added nodes of the tree.
- **toString, toTreeString:** The String representing a trie or sub-trie consists of a view of trie showing all the nodes, with added nodes differentiated from non-added, one node per line, and the binary tree relationships between the nodes visible, and with the element size (as described by `size`) of the trie and each sub-trie visible.

- **toAddedNodesTreeString**: A more compact representation of a trie (compared to `toString` or `toTreeString`) showing only the root and the added nodes, one node per line, and the non-binary tree relationships between the added nodes visible. This visualizes the tree produced by `constructAddedNodesTree`.

## Trie Iterators, Spliterators, and Streams

Most of these traversal methods traverse the nodes, while some traverse the keys/addresses of the nodes, including those iterators and spliterators within `AddressTrieSet` or `AddressTrieMap`.

Of those that traverse the nodes, some traverse only the added nodes and elements, while others traverse all the nodes. The added nodes are those explicitly added to the trie, the non-added nodes are those auto-generated for the binary trie structure. A non-added node can become added, and vice-versa.

Those that traverse the keys/addresses traverse only the added keys/addresses.

- **iterator, descendingIterator, nodeIterator, allNodeIterator, spliterator, descendingSpliterator, stream, parallelStream, nodeSpliterator, allNodeSpliterator**: Traverses the trie (or corresponding set) in natural trie order, either in forward (ascending) or reverse (descending) order. To obtain the corresponding stream for the two node spliterators, which do not have a corresponding `Stream` method, use `java.util.stream.StreamSupport.stream(Spliterator spliterator, boolean parallel)`.
- **containingFirstIterator, containingFirstNodeIterator, containingFirstAllNodeIterator**: Traverses the trie (or corresponding set) with a pre-order binary tree traversal. Such a traversal visits parent (containing prefix blocks) nodes before sub-nodes. Sub-nodes can be visited either lower sub-node first (`forwardSubNodeOrder`) or upper sub-node first. The node iterators allow you to cache an object with either sub-node when visiting the parent node, an object that can be retrieved when visiting the respective sub-node.
- **containedFirstIterator, containedFirstNodeIterator, containedFirstAllNodeIterator**: Traverses the trie (or corresponding set) with a post-order binary tree traversal. Such a traversal visits parent (containing prefix blocks) nodes after sub-nodes. Sub-nodes can be visited either lower sub-node first (`forwardSubNodeOrder`) or upper sub-node first.
- **blockSizeIterator, blockSizeNodeIterator, blockSizeAllNodeIterator, blockSizeCachingAllNodeIterator**: Traverses the trie from largest prefix blocks to smaller prefix blocks and then to individual addresses. The node iterators allow you to cache an object with either sub-node when visiting the parent node, an object that can be retrieved when visiting the respective sub-node.

See [the javadoc for TreeOps](#) for more details on the orderings for the various traversals.

## Parse String Representations of MAC Address

Parsing is like that for IP address. `MACAddressString` is used to convert. You can use one of `getAddress` or `toAddress`, the difference being whether parsing errors are handled by exception or not.

```
MACAddress address = new
MACAddressString("01:02:03:04:0a:0b").getAddress();
if(address != null) {
    //use address
}
```

or

```
try {
    MACAddress address = new MACAddressString("01:02:03:04:0a:0b").toAddress();
    //use address
} catch (AddressStringException e) {
    String msg = e.getMessage(); // detailed message indicating issue
}
```

## Various Formats of MAC Addresses

MAC Addresses are expected to be in hexadecimal. However, there is a number of accepted formats for MAC addresses:

```
aa:bb:cc:dd:ee:ff
aa-bb-cc-dd-ee-ff
aa bb cc dd ee ff
aabb.ccdd.eeff
aabbccddeeff
aabbcc-ddeeff
```

For the non-segmented format (aabbccddeeff), all 12 digits are required to avoid ambiguity.

MAC addresses can be either 48 or 64 bits, and so for each 48-bit format there is a 64-bit equivalent:

```
aa:bb:cc:dd:ee:ff:11:22
aa-bb-cc-dd-ee-ff-11-22
aa bb cc dd ee ff 11 22
aabb.ccdd.eeff.1122
aabbccddeeff1122
aabbcc-ddeeff1122
```

For the non-segmented 64-bit format (aabbccddeeff1122), all 16 digits are required to avoid ambiguity.

As with IP addresses, you can specify ranges using '\*' and '-' like aa-bb:\*:\*:cc:dd:ee. The range character for addresses that use the dash '-' character as a separator is '|', like aa|bb-\*-\*:cc-dd-ee.

## Format Examples

For instance, the address 0a:0b:0c:0d:0e:0f can be represented many ways:

```
static void parseMAC(String formats[]) {
    for(String format : formats) {
        System.out.println(new MACAddressString(format).getAddress());
    }
}

public static void main(String[] args) {
    String formats[] = {
        "a:b:c:d:e:f",
        "0a:0b:0c:0d:0e:0f",
        "a:b:c:d:e:f",
        "0a-0b-0c-0d-0e-0f",
        "0a0b0c-0d0e0f",
    }
```

```

        "0a0b.0c0d.0e0f",
        "0a 0b 0c 0d 0e 0f",
        "0a0b0c0d0e0f"
    };
    parseMAC(formats);
}

```

Output:

```

0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f
0a:0b:0c:0d:0e:0f

```

Ranges can be specified in the same way as for IP addresses ('-' for a specific range or '\*' for full range segments).

For instance, the address range ff:0f:aa-ff:00-ff:00-ff:00-ff can be represented many ways:

```

static void parseMAC(String formats[]) {
    for(String format : formats) {
        System.out.println(new MACAddressString(format).getAddress());
    }
}

public static void main(String[] args) {
    String formats[] = {
        "ff:f:aa-ff:00-ff:00-ff:00-ff",
        "ff:f:aa-ff:*:*:*",
        "ff:0f:aa-ff:*:*:*",
        "ff-0f-aa|ff-*-*-*",
        "ff0faa|ff0fff-*",
        "ff0f.aa00-ffff.*",
        "ff 0f aa-ff * * *",
        "ff0faa000000-ff0fffffffffff"
    };
    parseMAC(formats);
}

```

Output:

```

ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*
ff:0f:aa-ff:*:*:*

```

Delimited Segments

The range formats allow you to specify ranges of values. However, if you wish to parse addresses in which values are delimited, then you can use the methods `parseDelimitedSegments(String)` and `countDelimitedAddresses(String)` in `MACAddressString`. `parseDelimitedSegments` will provide an iterator to traverse through the individual addresses.

For example, given "1,2:3:4,5:6:7:8", `countDelimitedAddresses` will return 4 for the possible combinations: "1:3:4:6:7:8", "1:3:5:6:7:8", "2:3:4:6:7:8" and "2:3:5:6:7:8". With each string obtained from `parseDelimitedSegments` you can construct a `MACAddressString` instance.

## MAC Address Validation Options

Validation options allow you to restrict the allowed formats.

You have the class `MACAddressStringParameters` which can be passed into `MACAddressString`

You can restrict validation options using the appropriate builder classes that are nested classes in the options classes. Here is an example:

```
MACAddressStringParameters MAC_ADDRESS_OPTIONS_EXAMPLE =
    new MACAddressStringParameters.Builder().
        allowEmpty(false).
        allowAll(false).
        getFormatBuilder().
        setRangeOptions(RangeParameters.*NO_RANGE*).
        allowLeadingZeros(true).
        allowUnlimitedLeadingZeros(false).
        allowWildcardedSeparator(false).
        allowShortSegments(true).
        getParentBuilder().
        toParams();
```

The default options used by the library are permissive and not restrictive.

## MAC Address Prefix Lengths

A 'prefix length' in this library is defined generally as the bit-length of the foremost portion of the address that is not specific to an individual address but common amongst a group of addresses. For IP addresses, this is explicitly defined as part of the address using the '/' character. For IP addresses, the prefix is potentially variable and depends on how subnets have been allocated.

MAC Addresses don't have the exact same concept of prefix length as IP addresses. But concept of prefix can be applied in an implicit sense. For instance, a MAC address is typically divided into an OUI (Organizationally Unique Identifier) and ODI (Organizational Defined Identifier), so you might consider the OUI bits as the prefix. There are other ways of assigning MAC address blocks as well, such as IAB (individual address block), or MA-S/MA-M/MA-L (MAC Address block small, medium, and large), in which a certain number of higher bits are provided as an identifier to organizations from which they can create various extended identifiers using the lower bits. There is generally a pre-defined set of high bits that can be considered a prefix. This prefix is not variable and was typically assigned by the IEEE. However, there is no explicit way to represent a MAC address string with an associated prefix, as there is with CIDR IP addresses, so a prefix with MAC addresses is implicit rather than explicit.

Within this library, the prefix for a MAC address is defined as the largest number of high bits for which an address represents all addresses with the same set of higher bits.

For instance, the prefix length of `aa:::*` is 8 because the address represents all addresses that start with the same 8 bits "aa". The prefix length of `aa::cc:*` is 24 because the address represents all addresses that start with the same 16 bits `aa::cc`. The address `aa:bb:cc:dd:ee:ff` does not have a prefix or prefix length as it represents just a single address.

Once a MAC address has an associated prefix length, that prefix length remains the same while any operations are applied to the address (with the exception of operations that explicitly change the prefix, like `setPrefixLength`).

In summary, on the MAC side, the prefix length is implicit and based upon the address itself, while on the IP address side, the prefix length is explicitly defined.

## MAC Address Operations

---

Many of the same address operations available for IP addresses are available for MAC addresses, including the prefix operations, the section and segment access methods, iterators, spliterators, streams, containment, and reversal of bits, bytes and segments.

The reverse operations may be useful for "MSB format", "IBM format", "Token-Ring format", and "non-canonical form", where the bits are reversed in each byte of a MAC address.

## IPv6 – MAC Address Integration

---

There is a standardized procedure for converting MAC addresses to IPv6 given an IPv6 64-bit prefix, as described in [IETF RFC 2464](#), [RFC 3513](#), and [RFC 4944](#).

It details how to combine a 48 bit MAC address with a 64-bit IPv6 network prefix to produce an associated 128-bit IPv6 address. This is done by first constructing a 64-bit extended unique IPv6 interface identifier (EUI-64) from the MAC address. This library has implemented the same MAC / IPv6 integration.

Starting with a MAC address or section and with the IPv6 prefix, you can construct the associated IPv6 address with one of these constructors:

```
public IPv6Address(IPv6Address prefix, MACAddress eui)

public IPv6Address(IPv6AddressSection section, MACAddress eui)

public IPv6Address(IPv6AddressSection section, MACAddressSection
eui)

public IPv6Address(IPv6AddressSection section, MACAddressSection
eui, CharSequence zone)
```

There are equivalent methods in the class `MACAddress` for producing the link local address which has a pre-defined prefix, or for producing the host (interface identifier) address section of an IPv6 address.

```
public IPv6Address toLinkLocalIPv6()
public IPv6AddressSection toEUI64IPv6()
```

There is a similar method in `MACAddressSection`



```
public IPv6AddressSection toEUI64IPv6()
```

A MAC address is either 48 or 64 bits. To be converted to IPv6, the 48 bit address has segments inserted (two 1 byte MAC segments with value 0xfffe) to extend the address to 64 bits. For a 64 bit MAC address to be convertible, those same two segments must match the expected value of 0xffe. The following methods in MACAddress check will 64 bit addresses to ensure that are compatible with IPv6 by checking that the value of those two segments matches 0xfffe. Note that the asMAC argument allows you to extend using 0xffff rather than 0xfffe which is another manner by which a 48 bit MAC address can be extended to 64 bits. However, for purposes of extending to an IPv6 address the argument should be false so that the EUI-64 format is used.

```
public boolean isEUI64(boolean asMAC)
public MACAddress toEUI64(boolean asMAC)
```

There are similar methods in MACAddressSection

```
public boolean isEUI64(boolean asMAC)

public boolean isEUI64(boolean asMAC, boolean partial)

public MACAddressSection toEUI64(boolean asMAC)
```

Given an existing EUI-64 section you can use the prepend and append methods to create a full IPv6 address section containing all segments:

```
public IPv6AddressSection prepend(IPv6AddressSection other)

public IPv6AddressSection append(IPv6AddressSection other)
```

and from there you just construct the address:

```
public IPv6Address(IPv6AddressSection section)
```

To go the reverse direction IPv6 to MAC, there is a method in IPv6Address to produce a MAC address:

```
public MACAddress toEUI(boolean extended)
```

and another in IPv6AddressSection that uses whatever part of the interface identifier is included in the section to produce a MAC address section:

```
public MACAddressSection toEUI(boolean extended)
```

For instance, if the IPV6 address section is the 8 bytes corresponding to the network prefix of an IPV6 address section, then the resulting MAC address section will be 0 bytes. If the IPV6 address section is the 8 bytes corresponding to the interface identifier, and that identifier has the required 0xfffe values in the 5<sup>th</sup> and 6<sup>th</sup> bytes, then the MAC address section will be the full 8 bytes of an EUI-64 MAC address.

The following code is an example of constructing IPv6 addresses from a MAC address:



```

public static void main(String args[]) {
    MACAddressString macStr = new MACAddressString("aa:bb:cc:dd:ee:ff");
    MACAddress macAddress = macStr.getAddress();
    IPv6Address linkLocal = macAddress.toLinkLocalIPv6();
    System.out.println(linkLocal);

    IPAddressString ipv6Str = new
    IPAddressString("1111:2222:3333:4444::/64");
    IPv6Address ipv6Address = ipv6Str.getAddress().toIPv6();
    IPv6Address macIpv6 = new IPv6Address(ipv6Address, macAddress);
    System.out.println(macIpv6);
}

```

Output:

```

fe80::a8bb:ccff:fedd:eeff
1111:2222:3333:4444:a8bb:ccff:fedd:eeff/64

```

## Address Framework

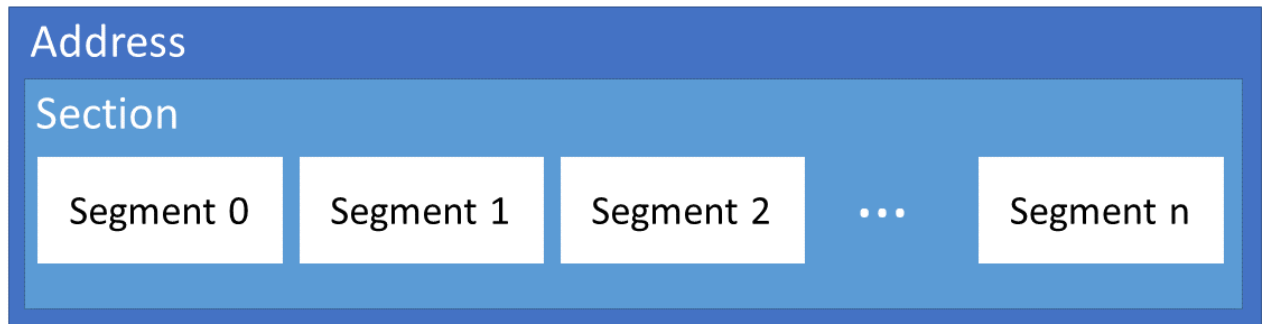
---

Much like there is a Java collections framework, there is an address framework to the IPAddress library. It is a unified set of inter-related interfaces, abstract implementations, and algorithms for all addresses and address components. It allows you to manipulate these items independently of implementation details, and provides standard interfaces to addresses and address components for code reuse and polymorphism. It represents the common structure of addresses, sections, segments, and so on.

You might wish to manipulate address components of different shapes and sizes transparently, or you may wish to manipulate different types of addresses or different address versions transparently. One element of the interface is the ability to convert any address component to bytes, whether division, segment, section, or address.

There is a hierarchy for the standard Address and Address Component data structures, which are addresses, sections of addresses, and segments of equal byte size inside those address sections.

Segments are equal length, each an exact number of bytes




Examples

IPv4: 0.0.255.255

IPv6: fe80:0:0:0:204:61ff:aaab:bbbb

MAC: aa-bb-cc-dd-ee-ff

Sections have a variable number of segments



There is a more diversified hierarchy for non-standard address structures, in which addresses or address sections might be divided into divisions of unequal length, or of non-integer byte-size.

Addresses can be split into sections or reconstituted into division groupings.  
Divisions need not be equal length, nor start and end on byte boundaries

## Division Grouping

Division 0

Division 1

Division 2

...


Division n

### Examples

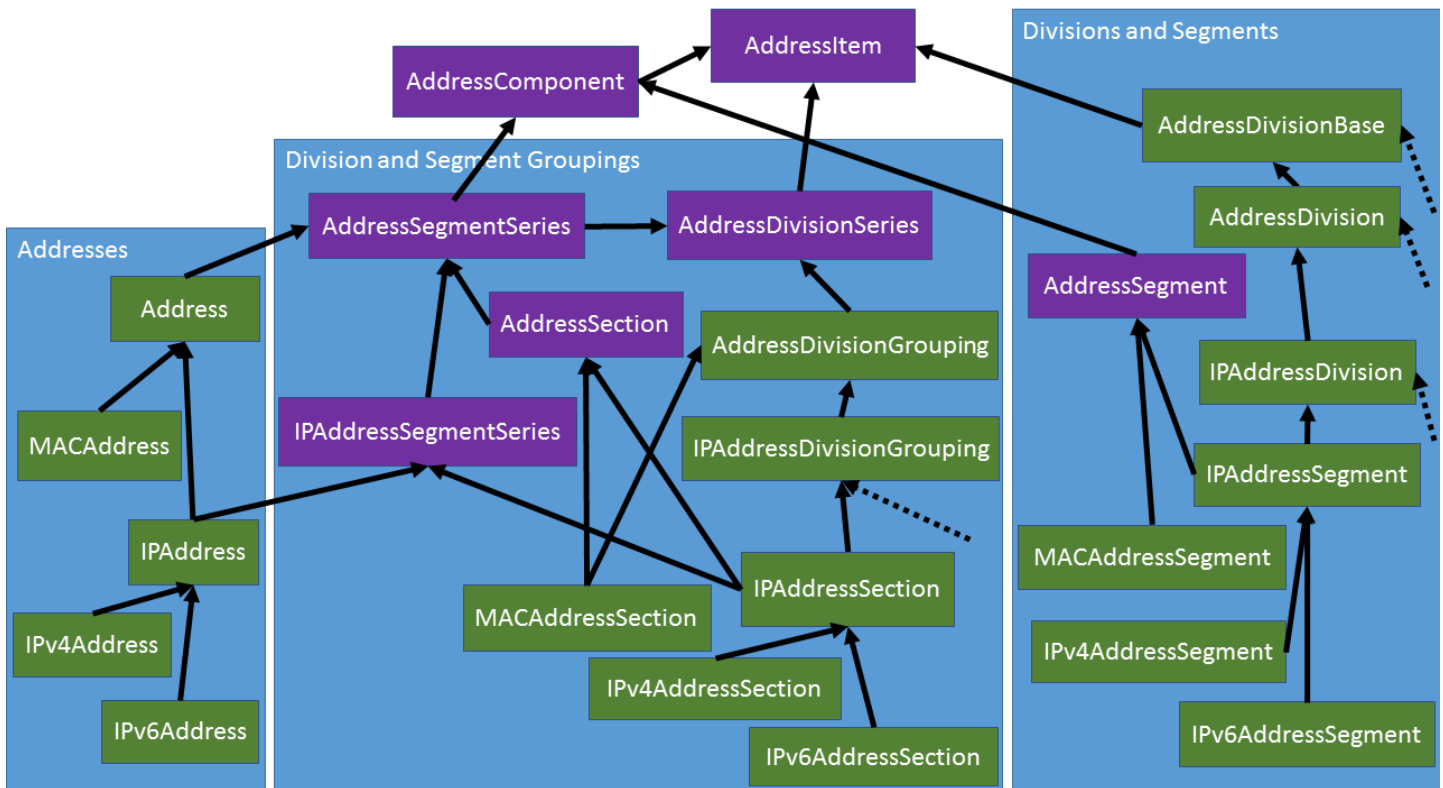
IPv4: 0.0.65535

IPv6 with IPv4 embedded: fe80:0:0:0:204:61ff:254.157.241.86

The address hierarchy of interfaces (purple) and classes (green) is shown:



Most of the full class hierarchy for address structure showing addresses, sections, division groupings, segments and divisions is shown here, separated into the three primary categories shown above. The dashed lines indicate there are a few less-prominent classes in the library not shown in the diagram.



## Conversion to String Representation of Address

Here is a list of string methods, in no specific order.

Note that, for a given address, the string produced by one of these methods might match those of other methods. They are not necessarily distinct from each other for all addresses.

All address types, HostName and IPAddressString:

- **toNormalizedString**: Produces a consistent string. For addresses, a string that is somewhat similar and consistent for all address components of the same type. Uses no segment compression in IPv6. Uses xx:xx:xx:xx:xx:xx for MAC address. Prints HostName instances in standard formats as expected by URLs and as dictated by RFCs, including port or service name and using square brackets for IPV6 addresses. Prints a standardized format for IPAddressString instances that cannot be converted to addresses (such as '\*'). For HostName and IPAddressString instances constructed from invalid strings, prints the original string used to construct.
- **toString**: same as **toCanonicalString** for addresses. For HostName and IPAddressString instances, prints the original string used to construct.

All address types:

- **toCanonicalString**: the string recommended by one or more RFCs or other common standard.
- **toCompressedString**: a short representation of the address while remaining within the confines of standard representation(s) of the address. For IPv6 compresses all compressible segments.
- **toHexString**: base 64 with no segments

IP addresses and HostName:

- **toNormalizedWildcardString**: similar to **toNormalizedString**, but uses wildcards and does not print prefix length for addresses that have prefix lengths

IP addresses only:

- **toFullString**: a string which maintains a common length, a string with no compressed segments and all segments of full length, which is 4 characters for IPv6 segments and 3 characters for IPv4 segments
- **toCanonicalWildcardString**: similar to **toCanonicalString**, but uses wildcards and does not print prefix length for addresses that have prefix lengths
- **toCompressedWildcardString**: similar to **toCompressedString**, but uses wildcards and does not print prefix length for addresses that have prefix lengths
- **toSQLWildcardString**: similar to **toNormalizedWildcardString**, but uses SQL wildcards
- **toPrefixLengthString**: a string with CIDR prefix length if it has one, and compresses the host for IPv6. For IPv4 it is the same as **toCanonicalString**.
- **toSubnetString**: uses **toNormalizedWildcardString** for IPv4 and **toPrefixLengthString** for IPv6
- **toReverseDNSLookupString**: a string for reverse DNS lookup
- **toOctalString**: base 8, optionally with a '0' prefix to indicate octal
- **toBinaryString**: all ones and zeros
- **toConvertedString**: if an IPv6 address can be converted to IPv4 as determined by `isIPv4Convertible()`, produces the result of **toMixedString**. For IPv4 same as **toCanonicalString**.
- **toUNCHostName**: Microsoft UNC path component
- **toNormalizedString(IPStringOptions)** : your own customized string

IPv6 only:

- **toMixedString**: mixed IPv6/IPv4 string format like a:b:c:d:e:f:255.0.255.255
- **toBase85String**: RFC 1924

MAC address only:

- **toColonDelimitedString**: same as **toNormalizedString**, xx:xx:xx:xx:xx:xx
- **toDashedString**: same as **toCanonicalString**, xx-xx-xx-xx-xx-xx
- **toDottedString**: xxxx.xxxx.xxxx
- **toSpaceDelimitedString**: xx xx xx xx xx xx

## More Details and Examples

`AddressSegmentSeries` objects (such as `Address` and `AddressSection`) have methods that produce strings in various formats: `toCanonicalString`, `toNormalizedString`, `toCompressedString`, and `toHexString`.

```
public static void main(String[] args) {
    printStrings(new MACAddressString("a:bb:c:dd:e:ff").getAddress());
}
```

```

printStrings(new MACAddressString("a:bb:c:*").getAddress());
printStrings(new IPAddressString("a:bb:c::/64").getAddress());
printStrings(new IPAddressString("1.2.3.4").getAddress());
printStrings(new IPAddressString("1.2.0.0/16").getAddress());
printStrings(new IPAddressString("a:bb:c::dd:e:ff").getAddress());
}

static void printStrings(AddressSegmentSeries series) {
    System.out.println(series.toCanonicalString());
    System.out.println(series.toNormalizedString());
    System.out.println(series.toCompressedString());
    System.out.println(series.toHexString(true));
    System.out.print("lower: " + series.getLower() + " bytes:");
    System.out.println(Arrays.toString(series.getBytes()));
    System.out.print("upper: " + series.getUpper() + " bytes:");
    System.out.println(Arrays.toString(series.getUpperBytes()));
    System.out.println();
}

```

Output:

```

0a-bb-0c-dd-0e-ff
0a:bb:0c:dd:0e:ff
a:bb:c:dd:e:ff
0x0abb0cdd0eff
lower: 0a:bb:0c:dd:0e:ff bytes:[10, -69, 12, -35, 14, -1]
upper: 0a:bb:0c:dd:0e:ff bytes:[10, -69, 12, -35, 14, -1]

0a-bb-0c-*-*-*
0a:bb:0c:*:*:~*
a:bb:c:*:*:~*
0x0abb0c000000-0x0abb0cffffff
lower: 0a:bb:0c:00:00:00 bytes:[10, -69, 12, 0, 0, 0]
upper: 0a:bb:0c:ff:ff:ff bytes:[10, -69, 12, -1, -1, -1]

a:bb:c::/64
a:bb:c:0:0:0:0:0/64
a:bb:c::/64
0x000a00bb000c00000000000000000000-0x000a00bb000c0000ffffffffffffffff
lower: a:bb:c:0:0:0:0:0 bytes:[0, 10, 0, -69, 0, 12, 0, 0, 0, 0, 0, 0, 0, 0, 0]
upper: a:bb:c:0:ffff:ffff:ffff:ffff bytes:[0, 10, 0, -69, 0, 12, 0, 0, -1, -1, -1, -1, -1, -1, -1]

1.2.3.4
1.2.3.4
1.2.3.4
0x01020304
lower: 1.2.3.4 bytes:[1, 2, 3, 4]
upper: 1.2.3.4 bytes:[1, 2, 3, 4]

1.2.0.0/16
1.2.0.0/16
1.2.0.0/16
0x01020000-0x0102ffff
lower: 1.2.0.0 bytes:[1, 2, 0, 0]
upper: 1.2.255.255 bytes:[1, 2, -1, -1]

a:bb:c::dd:e:ff
a:bb:c:0:0:dd:e:ff
a:bb:c::dd:e:ff
0x000a00bb000c0000000000dd000e00ff
lower: a:bb:c:0:0:dd:e:ff bytes:[0, 10, 0, -69, 0, 12, 0, 0, 0, 0, 0, 0, 0, 0, 0]

```

```
-35, 0, 14, 0, -1]
upper: a:bb:c:0:0:dd:e:ff bytes:[0, 10, 0, -69, 0, 12, 0, 0, 0, 0, 0,
-35, 0, 14, 0, -1]
```

The `IPAddress` and `IPAddressSection` classes and their version-specific subclasses have additional methods to produce specific strings representing the address:

```
public static void main(String[] args) {
    IPAddress address = new IPAddressString("a:b:c::e:f").getAddress();
    print(address);
    address = new IPAddressString("a:b:c::").getAddress();
    print(address);

    //a:b:c:*:: cannot be represented as a range of two single values
    //so this throws exception in toBase85String(), toBinaryString(), and toHexString()
    address = new IPAddressString("a:b:c:*::").getAddress();
    print(address);
}

static void print(IPAddress address) {
    System.out.println(address.toCanonicalString());
    System.out.println(address.toFullString());
    System.out.println(address.toNormalizedString());
    System.out.println(address.toSQLWildcardString());
    System.out.println(address.toSubnetString());
    try {
        if(address.isIPv6()) {
            System.out.println(address.toIPv6().toMixedString());
            System.out.println(address.toIPv6().toBase85String());
        }
        System.out.println(address.toBinaryString());
        System.out.println(address.toHexString(true));
    } catch(IncompatibleAddressException e) {}
    System.out.println();
}
```

Output:

[illegible]





```
0.0.0.0.0.0.0.0.0.8.b.d.0.1.0.0.2.ip6.arpa
*.*.*.*.*.*.*.*.*.*.*.*.*.*.0.0.0.0.0.0.0.0.0.8.b.d.0.1.0.0.2.ip6.arpa
0.0.0.0.0.0.0.0.0.8.b.d.0.1.0.0.2.ip6.arpa
```

## General String Methods

The methods `toCanonicalString` and `toCompressedString` are available for any address or address section.

The methods `toHexString` and `toNormalizedString` are available for any address component, including segments and sections.

## Prefix Length Indicator in Strings

Typically prefix lengths will be added to IP strings. To choose to print wildcards \* and range characters – as opposed to using prefix length, there are additional methods:

```
public static void main(String[] args) {
    IPAddress address = new IPAddressString("a:b:c::/64").getAddress();
    print(address);
    address = new IPAddressString("a:b:c*::/64").getAddress();
    print(address);
}

static void print(IPAddress address) {
    System.out.println(address.toCanonicalString());
    System.out.println(address.toCanonicalWildcardString());
    System.out.println(address.toFullString());
    System.out.println(address.toNormalizedString());
    System.out.println(address.toSQLWildcardString());
    System.out.println();
}
```

Output:

```
a:b:c::/64
a:b:c:0::*:*:~*~*~*
000a:000b:000c:0000:0000:0000:0000:0000/64
a:b:c:0:0:0:0:0/64
a:b:c:0::*:*:~*~*~*
a:b:c::*/64
a:b:c::*:*:~*~*~*
000a:000b:000c:0000-ffff:0000:0000:0000:0000/64
a:b:c::0:0:0:0:0/64
a:b:c::*:*:~*~*~*
```

## IP Version-dependent Strings

Some strings are version-dependent:

```
public static void main(String[] args) {
    // toSubnetString() prefers '*' for IPv4 and prefix for IPv6
    System.out.println(new IPAddressString("1.2.0.0/16").getAddress().toSubnetString());
}
```

```

System.out.println(new IPAddressString("a:b::/64").getAddress().toSubnetString());

//converts IPv4-mapped to a:b:c:d:e:f:1.2.3.4 notation
System.out.println(new IPAddressString("::ffff:a:b").getAddress().toConvertedString());
}

```

Output:

```

1.2.*.*
a:b::/64
::ffff:0.10.0.11

```

## Collections of IP Address Strings

You can produce collections of strings:

```

public static void main(String[] args) {
    IPAddress address = new IPAddressString("a:b:c::e:f").getAddress();
    print(address);
}

private static void print(IPAddress address) {
    //print(address.toAllStrings()); produces many strings
    print(address.toStandardStrings());
}

public static void print(String strings[]) {
    for(String str : strings) {
        System.out.println(str);
    }
}

```

Output:

```

a:b:c:0:0:0:0.14.0.15
a:b:c:0:0:0:000.014.000.015
000a:000b:000c:0000:0000:0000:0.14.0.15
000a:000b:000c:0000:0000:0000:000.014.000.015
A:B:C:0:0:0:0.14.0.15
A:B:C:0:0:0:000.014.000.015
000A:000B:000C:0000:0000:0000:0.14.0.15
000A:000B:000C:0000:0000:0000:000.014.000.015
a:b:c::0.14.0.15
a:b:c::000.014.000.015
000a:000b:000c::0.14.0.15
000a:000b:000c::000.014.000.015
A:B:C::0.14.0.15
A:B:C::000.014.000.015
000A:000B:000C::0.14.0.15
000A:000B:000C::000.014.000.015
a:b:c:0:0:0:e:f
000a:000b:000c:0000:0000:0000:000e:000f
A:B:C:0:0:0:E:F
000A:000B:000C:0000:0000:0000:000E:000F
a:b:c::e:f
000a:000b:000c::000e:000f

```

```
A:B:C::E:F
000A:000B:000C::000E:000F
```

The String collections can be customized with `toNormalizedString(StringOptions params)`

Note that string collections never have duplicate strings. The String collections can be customized with `toStrings(IPStringBuilderOptions options)` .

## Searching Text of Databases for all Addresses in a Subnet

---

Suppose you wanted to search for all addresses from a subnet in a large amount of text data. For instance, suppose you wanted to search for all addresses in the text from the subnet a:b:0:0::/64. You can start a representation of just the network prefix section of the address, then you can get all such strings for that prefix.

```
IPAddressSection prefix = new IPAddressString("a:b::").getAddress().
    getNetworkSection(64, false);
String strings[] = prefix.toStandardStringCollection().toStrings();
for(String str : strings) {
    System.out.println(str);
}
```

Output:

```
a:b:0:0
000a:000b:0000:0000
A:B:0:0
000A:000B:0000:0000
a:b::
000a:000b::
A:B::
000A:000B::
```

If you need to be more stringent or less stringent about the address formats you wish to search, then you can use `toStringCollection(IPStringBuilderOptions options)` with an instance of `IPv6StringBuilderOptions` .

Searching for those strings will find the subnet addresses. However, you may get a few false positives, like "a:b::d:e:f:a:b". To eliminate the false positives, you can just emulate in Java the SQL code produced below for the SQL database search, using substrings constructed from the segment separators.

For a MySQL database search:

```
public static void main(String[] args) {
    IPAddressSection prefix = new IPAddressString("a:b::").
        getAddress().getNetworkSection(64, false);
    StringBuilder sql = new StringBuilder("Select rows from table where");
    prefix.getStartsWithSQLClause(sql, "column1");
    System.out.println(sql);
}
```

Output:

```
Select rows from table where ((substring_index(column1,':',4) =
'a:b:0:0') OR ((substring_index(column1,':',3) = 'a:b:') AND (LENGTH
(column1) - LENGTH(REPLACE(column1, ':', '')) <= 6)))
```

For IPv4, another way to search for a subnet like 1.2.0.0/16 would be to do an SQL `SELECT` with the SQL wildcard string:

```
public static void main(String[] args) {
    String wildcardString = new IPAddressString("1.2.0.0/16").
        getAddress().toSQLWildcardString();
    System.out.println(wildcardString);
}
```

Output:

```
1.2.%.%
```

Then your SQL search string would be like:

```
Select rows from table where column1 like 1.2.%.%
```

## Containment and Subnet Membership

---

To check whether an IP address is contained by a subnet:

```
IPAddress address = new
IPAddressString("1.2.0.0/16").getAddress();
System.out.println(address.contains(new
IPAddressString("1.2.3.4").getAddress()));
System.out.println(address.contains(new
IPAddressString("1.2.3.0/24").getAddress()));
System.out.println(address.contains(new
IPAddressString("1.2.3.0/25").getAddress()));
System.out.println(address.contains(new
IPAddressString("1.1.0.0").getAddress()));
```

Output:

```
true
true
true
false
```

The `contains` method is not restricted to IP addresses or IP address prefixed addresses. There is a `contains` method for every `Address` or `AddressSection`.

For checking containment of an address or subnet in a large number of subnets, or to check containment of a large number of addresses or subnets in a subnet, use the address trie data structure, using one of the subclasses of `AddressTrie`.

There is also an assortment of iterators for addresses, sections, and segments which represent multiple values. There is an `iterator()`, `getLower()` method and `getUpper()` method for every address component.

## DNS Resolution and URLs

---

If you have a string that can be a host or an address and you wish to resolve to an address, create a `HostName` and use `HostName.toResolvedAddress()`. If you wish to obtain a string representation to be part of a URL, use `HostName.toNormalizedString()`.

## Sorting and Comparisons

---

Comparing and sorting can be useful for storing addresses in certain types of data structures. All of the core classes implement `java.lang.Comparable`. In fact, any `AddressItem` is comparable to any other, which covers almost every type in the address framework. Different representations of the same address or subnet are considered equal. Different representations of the same set of addresses are considered equal. However, `HostName` instances are not equal to `IPAddressString` instances nor `IPAddress` instances, and instances of `IPAddressString` are not equal to instances of `IPAddress`, not even when representing the same address or subnet.

The library provides the abstract type `AddressComparator` and some implementations for comparison purposes. Address classes use the subclass `CountComparator` for their natural ordering. When comparing subnets, you can either emphasize the count of addresses, or you can emphasize the values of the lower or upper address represented by the subnet, and comparators are provided for those variations.

The address tries provide their own sorting and comparison, which matches the comparators above with respect to individual addresses, but when comparing a prefix block to a second address or block with larger prefix, and the first prefix matches the same bits in the second, the trie comparator orders by the bit that follows the first prefix in the second.

## Cache Classes

---

The `IPAddressNetwork` class defines `IPAddressStringCache` and `HostNameCache`. These cache classes allow you to cache identifier strings (hosts and/or addresses). Note that the identifier strings themselves will also cache their associated addresses, whether parsed or resolved. Also note that addresses cache their associated strings and various other objects. Therefore, these cache classes go a long way towards allowing you to avoid creating the same objects frequently. These caches do quick lookups using either bytes or strings, which can be ideal for some applications that handle many addresses or host names.

## Make your IPv4 App work with IPv6

---

If you want to make your app work with IPv6, that is a wise decision. IPv6 penetration is getting close to 50% in some countries.

But you have no IPv6 testing and a lot of code.

Start by replacing your IPv4 code with the types and operations in this library. But don't use the IPv4-specific types like `IPv4Address` everywhere, use `IPAddressString`, `IPAddress`, and the other polymorphic types, along with their polymorphic operations such as `toString`, `contains`, and `iterator` that avoid exposing the specifics of IPv4. If your app

uses network or host masks, start working with CIDR prefix lengths instead. Use methods like [getBlockMaskPrefixLength](#) to store prefix lengths, even if the app continues to use masking.

When you're ready, use your existing regression tests, ensuring the app still runs as always on IPv4. Once your existing IPv4 tests all pass, you are already most of the way to supporting IPv6, without having written a single line of IPv6 code. With some apps, you may even be ready to try it out with IPv6.

## Alternative Options for Prefix Length Handling

This feature is provided primarily for backwards compatibility with versions 3 and earlier.

As indicated in the section on parsing, when constructing an IP address with prefix length, you can end up with an individual address or a subnet if you provide a network address. This is also true for address and address sections constructed from numeric values rather than strings. When you specify the network address of a subnet with its corresponding prefix length, then the resulting object will be the subnet, and not an individual address.

In versions 3 and earlier, prefixed addresses and sections were always converted to the subnet block of addresses for that prefix, regardless of whether you specified the network address or any other address with a prefix length.

In the current version, there are three options available, the version 3 legacy behavior, the new default behavior, and a third option:

1. **Legacy behavior:** An address with a prefix length is always the subnet block for that prefix. This is the legacy behavior, the behaviour for versions 3 and under of this library. An address like 1.2.3.4/16 is equivalent to 1.2.\*/16, and a:b:c:d::/64 is equivalent to a:b:c:d:\*/64.
2. **New default behavior:** An address with a prefix length and a zero-valued host is the subnet block for that prefix. An address with a host value of non-zero is a single address corresponding to a single host. This is the new default behavior for this library since it is a convention with common usage.

For example, the IPv4 address 1.2.3.4/16 is just a single address and not a subnet, while 1.2.0.0/16 is the prefix block, it is the subnet 1.2.\*/16.

This behavior is akin to the common convention used by many network administrators, many routers and many applications. With IPv4, an address with a host value of zero is known as the network address. It often represents the subnet of all addresses with that same prefix. It is often used for routing. The all-zero address 0.0.0.0 is conventionally known as `INADDR_ANY` (any address on the local machine), it is the address that returns true for `java.net.Inet4Address.isAnyLocalAddress()`, and when paired with prefix zero it is known as the default route (the route for all addresses).

With IPv6, the convention is similar and this option results in the same behaviour. It is a standard convention with IPv6 to denote a subnet like 1:2:3:4::/64, with a zero-valued (and usually compressed) host, and that subnet represents the prefix block all addresses starting with that prefix 1:2:3:4. An address with a host of zero is known as the anycast address. The all-zero address `::` is the value of `IN6ADDR_ANY_INIT`, the analog to the IPv4 `INADDR_ANY`, and the address that returns true for `java.net.Inet6Address.isAnyLocalAddress()`.

3. **Explicitly defined subnet:** The third option is the one in which a subnet block derived from a prefix length must be explicitly defined, you must list the range of values explicitly. Any subnet must be explicitly defined, like 1.2.\*/16, while 1.2.0.0/16 is just a single address.

Version 3 and earlier versions of the library provided only the first behavior. All three options are available starting with version 4 of this library, while the default behavior is the second options listed and is most consistent with common

conventions used widely for IP addresses.

The options above are generally geared towards IP addresses. With MAC Addresses prefix lengths are implicit for the most part (and cannot be specified as part of an address string explicitly). MAC address prefix lengths are discussed in more detail in the section below “MAC Address Prefix Lengths” With MAC there is no difference between options 2 and 3 since there is no analog to the zero-host/network/anycast address. With MAC, when using an operation that allows you to supply a prefix, the first option above results in addresses that span all values beyond the prefix, while options 2 and 3 do not.

## Configuring Alternative Prefix Length Handling

To choose one of the non-default prefix length handling options, such as the legacy option, you make use of the `PrefixConfiguration` enum type and the network objects.

Each address version/type (IPv4, IPv6, MAC) has an associated network object and each network object has a `PrefixConfiguration` setting. Prefix length has a similar meaning for MAC but with some key differences outlined below. By default, all networks use the same setting of `PrefixConfiguration.DEFAULT` stored in a single static field in the `AddressNetwork` class.

By calling `setPrefixConfiguration` in `AddressNetwork` you can change this default setting for all networks. Or, you can call `setPrefixConfiguration` in one or more of `IPv4AddressNetwork`, `IPv6AddressNetwork`, or `MACAddressNetwork`, to set each one individually. Typically, you would change the setting when the application starts. Should you choose to do so at a different point in the application lifetime, it should be noted that the behavior is undefined when using pre-existing address objects following such a change. It is recommended that the prefix length configuration remain constant throughout the lifetime of an application, or that existing address objects be discarded when the configuration is changed.

If discarding existing address objects, keep in mind that the library caches address and address component objects. To discard those cached objects you can clear those caches by calling `defaultIPv6Network()`, `defaultIPv4Network()`, and `defaultMACNetwork()` to get the default network objects and then call `clearCaches()` on each one.

If two applications using the `IPAddress` library at the same time in the same process require conflicting prefix configurations, then one option to achieve this is to use different class loaders in the separate apps, loading the `IPAddress` library in separate class loaders. In fact, many popular java applications already use class loaders for code separation, such as the Eclipse framework which uses them to separate plugins and the Apache Tomcat web server which uses them to separate web-apps.

If you cannot load the library in separate class loaders, then another option is available to you. This option is to use your own network classes. The procedure for using your own network classes is described above in the section of this document on networks. The first step is to subclass the address and address component classes. The next step is to override the key methods for accessing the networks in those classes to supply your own network subclasses which will have their own non-default configuration. When using your own network classes, you can override the `getPrefixConfiguration()` method of `AddressNetwork` to avoid using the default configuration setting and use your own. In your network, you will also need to override either the `createAddressCreator()` or `getAddressCreator()` methods to supply your own address creator subclass instance that will create your own address and address component classes.

The following example shows how to create an `IPAddressString` that will use a customized network class along with customized address component classes:

```
static class MyIPv6Address extends IPv6Address {

    public MyIPv6Address(byte[] bytes, Integer prefixLength) {
        super(bytes, prefixLength);
    }
}
```



```

    public MyIPv6Address(MyIPv6AddressSection section, CharSequence zone) {
        super(section, zone);
    }

    public MyIPv6Address(MyIPv6AddressSection section) {
        super(section);
    }

    @Override
    public IPv6AddressNetwork getNetwork() {
        return myIPv6Network;
    }
}

static class MyIPv6AddressSection extends IPv6AddressSection {

    public MyIPv6AddressSection(byte[] bytes, Integer prefixLength) {
        super(bytes, prefixLength);
    }

    public MyIPv6AddressSection(IPv6AddressSegment[] segments,
        int startIndex, boolean cloneSegments) {
        super(segments, startIndex, cloneSegments);
    }

    public MyIPv6AddressSection(IPv6AddressSegment[] segments,
        Integer prefixLength) {
        super(segments, prefixLength);
    }

    @Override
    public IPv6AddressNetwork getNetwork() {
        return myIPv6Network;
    }
}

static class MyIPv6AddressSegment extends IPv6AddressSegment {

    public MyIPv6AddressSegment(int lower,
        Integer segmentPrefixLength) {
        super(lower, segmentPrefixLength);
    }

    public MyIPv6AddressSegment(int lower, int upper,
        Integer segmentPrefixLength) {
        super(lower, upper, segmentPrefixLength);
    }

    @Override
    public IPv6AddressNetwork getNetwork() {
        return myIPv6Network;
    }
}

static IPv6AddressNetwork myIPv6Network = new IPv6AddressNetwork() {
    @Override
    public PrefixConfiguration getPrefixConfiguration() {
        return PrefixConfiguration.ALL_PREFIXED_ADDRESSES_ARE_SUBNETS;
    }

    @Override
    protected IPv6AddressCreator createAddressCreator() {
        return new IPv6AddressCreator() {
            @Override

```



```

public IPv6AddressSection createSection(byte bytes[],
    Integer prefix) {
    return new MyIPv6AddressSection(bytes, prefix);
}

@Override
public IPv6AddressSegment createSegment(int value,
    Integer segmentPrefixLength) {
    return new MyIPv6AddressSegment(value, segmentPrefixLength);
}

@Override
public IPv6AddressSegment createSegment(int lower,
    int upper, Integer segmentPrefixLength) {
    return new MyIPv6AddressSegment(lower, upper,
        segmentPrefixLength);
}

@Override
protected IPv6AddressSegment createSegmentInternal(int value,
    Integer segmentPrefixLength, CharSequence addressStr,
    int originalVal, boolean isStandardString,
    int lowerStringStartIndex, int lowerStringEndIndex) {
    return new MyIPv6AddressSegment(value, segmentPrefixLength);
}

@Override
protected IPv6AddressSection
    createPrefixedSectionInternal(IPv6AddressSegment segments[],
        Integer prefix) {
    return new MyIPv6AddressSection(segments, prefix);
}

@Override
protected IPv6AddressSection
    createSectionInternal(IPv6AddressSegment segments[]) {
    return new MyIPv6AddressSection(segments, 0, false);
}

@Override
protected IPv6Address createAddressInternal(
    IPv6AddressSection section, HostIdentifierString from) {
    return new MyIPv6Address((MyIPv6AddressSection) section);
}

@Override
protected IPv6Address createAddressInternal(
    IPv6AddressSection section, CharSequence zone,
    HostIdentifierString from) {
    return new MyIPv6Address((MyIPv6AddressSection) section, zone);
}
}
};

IPAddressStringParameters params = new IPAddressStringParameters.Builder().
    getIPv6AddressParametersBuilder().setNetwork(myIPv6Network).getParentBuilder().toParams();
IPv6Address myAddr = new IPAddressString("1::1/64", params).getAddress().toIPv6();

```